

Unsupervised Learning:

HR Employee Clustering Analysis

1. Introduction

This analysis explores an HR dataset containing employee demographic and job-related information such as gender, department, job title, location, hire dates, and employment status. The dataset, sourced from a publicly available Human Resources repository, includes over 22,000 valid employee records after cleaning.

The primary objective of this study was to apply unsupervised learning techniques to uncover hidden patterns among employees and group them into meaningful clusters. By performing clustering and dimensionality reduction, the analysis aims to help HR teams identify employee segments that can inform decision-making in retention, training, workforce planning, and engagement strategies.

2. Dataset Source and Description

The dataset used for this analysis was obtained from Figshare, a public research data repository that provides open access to academic datasets. The specific file, titled "Human Resources.csv," contains anonymized employee data designed for educational and analytical use in workforce analytics.

https://figshare.com/articles/dataset/Human_Resources_csv/28780886

The dataset contains approximately 37,000 employee entries, of which about 22,000 are complete and valid for modeling. It includes 13 columns describing demographic and employment information, such as ID, gender, race, department, job title, location (city and state), hire date, and term date. During preprocessing, additional numerical features — Age and Years of Service — were engineered from date fields to represent experience and tenure. The dataset's mix of categorical and temporal features makes it well suited for clustering and dimensionality reduction to reveal underlying workforce patterns.

3. Methodology

Data Preprocessing

Missing values were handled using median imputation for date fields and mode replacement for categorical variables. Irrelevant identifiers (e.g., names, employee IDs) were dropped to maintain anonymity and reduce noise. Categorical features such as department, job title, gender, race, and location were encoded using one-hot encoding

to enable numerical computation. All numerical variables were standardized using StandardScaler to ensure equal weight during clustering.

Clustering and Dimensionality Reduction

The K-Means algorithm was implemented to identify employee clusters based on shared characteristics. The Elbow Method and Silhouette Score were applied to determine the optimal number of clusters. Principal Component Analysis (PCA) reduced dimensionality and allowed for two-dimensional visualization of the cluster structures.

4. Results and Interpretation

The Elbow Method displayed a clear bend around $k = 4$, suggesting four natural clusters. Silhouette analysis produced scores of $k=2$ (0.456), $k=3$ (0.373), $k=4$ (0.322), $k=5$ (0.333), and $k=6$ (0.341). Although the highest score occurred at $k = 2$, the four-cluster solution ($k = 4$) was selected for its richer segmentation and better interpretability.

PCA visualization confirmed moderate separation among clusters. The groups likely represent:

New or early-career hires – low age and tenure

Mid-career employees – moderate age and years of service

Long-tenured, senior staff – high experience

Remote or specialized employees – distinctive locations or job roles

These segments reflect real organizational structures where roles overlap but exhibit distinct characteristics useful for HR analysis.

5. Model Evaluation

The final model's Silhouette Score of 0.322 indicates moderate cohesion and separation, common in HR datasets where workforce traits often overlap. The results confirm that the four-cluster model balances statistical validity with business interpretability. It provides actionable segmentation for HR teams to understand workforce composition and plan targeted initiatives such as professional development, retention programs, and onboarding support.

6. Challenges Encountered

Several challenges were encountered during the analysis. The dataset contained missing and inconsistent date values, which required median imputation to maintain data completeness. Managing categorical variables with many unique levels, such as job titles and cities, increased dimensionality after one-hot encoding, making model computation more complex. In addition, interpreting moderate Silhouette Scores was challenging

because employee attributes naturally overlap across departments and tenure levels, limiting perfect cluster separation. Despite these obstacles, careful preprocessing, feature engineering, and model evaluation ensured that the final clustering results remained both statistically reliable and practically interpretable.

7. Deployment and Monitoring Strategy

For deployment, the clustering model can be integrated into an HR analytics dashboard that updates automatically as new employee records are added. Employees can be assigned to clusters in real time to monitor workforce changes and identify emerging patterns.

Ongoing monitoring should include:

Quarterly retraining of the model to capture new hiring or attrition trends

Visualization dashboards to track cluster shifts over time

Feedback loops from HR managers to validate cluster interpretations

By embedding the model into HR systems, organizations can leverage data-driven segmentation to enhance decision-making in areas such as retention planning, workforce diversity, and employee engagement.

Double-click (or enter) to edit

```
1 # Import Various Libraries
2
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.model_selection import train_test_split
8 from sklearn.decomposition import PCA
9 from sklearn.cluster import KMeans
10 from sklearn.metrics import silhouette_score
11 import warnings
12
13 warnings.filterwarnings('ignore') # don't show me warnings going forward
14
15
16 # Load Dataset
17
18 hr_df = pd.read_csv('Human Resources.csv')
19
20 print('First 5 Heads Dataset')
21 print(hr_df.head())
22 print("-----")
```

```

24 # # Convert to Dataframe
25 print('Convert to DataFrame')
26 hr_df = pd.DataFrame(hr_df)
27 print(hr_df)
28 print("-----")
29
30
31

```

First 5 Heads Dataset

	id	first_name	last_name	birthdate	gender	\
0	00-0037846	Kimmy	Walczynski	6/4/1991	Male	
1	00-0041533	Ignatius	Springett	6/29/1984	Male	
2	00-0045747	Corbie	Bittlestone	7/29/1989	Male	
3	00-0055274	Baxy	Matton	9/14/1982	Female	
4	00-0076100	Terrell	Suff	4/11/1994	Female	

	race	department	\
0	Hispanic or Latino	Engineering	
1	White	Business Development	
2	Black or African American	Sales	
3	White	Services	
4	Two or More Races	Product Management	

	jobtitle	location	hire_date	\
0	Programmer Analyst I	Headquarters	1/20/2002	
1	Business Analyst	Headquarters	4/8/2019	
2	Solutions Engineer Manager	Headquarters	10/12/2010	
3	Service Tech	Headquarters	4/10/2005	
4	Business Analyst	Remote	9/29/2010	

	termdate	location_city	location_state	
0	NaN	Cleveland	Ohio	
1	NaN	Cleveland	Ohio	
2	NaN	Cleveland	Ohio	
3	NaN	Cleveland	Ohio	
4	2029-10-29 06:09:38 UTC	Flint	Michigan	

Convert to DataFrame

	id	first_name	last_name	birthdate	gender	\
0	00-0037846	Kimmy	Walczynski	6/4/1991	Male	
1	00-0041533	Ignatius	Springett	6/29/1984	Male	
2	00-0045747	Corbie	Bittlestone	7/29/1989	Male	
3	00-0055274	Baxy	Matton	9/14/1982	Female	
4	00-0076100	Terrell	Suff	4/11/1994	Female	
...
37403	NaN	NaN	NaN	NaN	NaN	NaN
37404	NaN	NaN	NaN	NaN	NaN	NaN
37405	NaN	NaN	NaN	NaN	NaN	NaN
37406	NaN	NaN	NaN	NaN	NaN	NaN
37407	NaN	NaN	NaN	NaN	NaN	NaN

	race	department	\
0	Hispanic or Latino	Engineering	
1	White	Business Development	

```

2      Black or African American           Sales
3                  White                 Services
4      Two or More Races   Product Management
...
37403                   ...                   ...
37404                   NaN                  NaN
37405                   NaN                  NaN
37406                   NaN                  NaN
37407                   NaN                  NaN

```

jobtitle location hire_date \

```

1 # # Inspect the Dataset:
2
3 print('Inspect the Dataset')
4 print(hr_df.info())
5 print("-----")
6
7 print('Describe the Dataset')
8 print(hr_df.describe())
9 print("-----")
10
11 print('Dataset Shape')
12 print(hr_df.shape)
13 print("-----")
14

```

Inspect the Dataset

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 37408 entries, 0 to 37407
Data columns (total 13 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                22214 non-null   object 
 1   first_name        22214 non-null   object 
 2   last_name         22214 non-null   object 
 3   birthdate         22214 non-null   object 
 4   gender             22214 non-null   object 
 5   race               22214 non-null   object 
 6   department         22214 non-null   object 
 7   jobtitle          22267 non-null   object 
 8   location           22214 non-null   object 
 9   hire_date          22214 non-null   object 
 10  termdate           3929 non-null   object 
 11  location_city     22214 non-null   object 
 12  location_state    22214 non-null   object 
dtypes: object(13)
memory usage: 3.7+ MB
None
-----
```

Describe the Dataset

	id	first_name	last_name	birthdate	gender	race	department	\
count	22214	22214	22214	22214	22214	22214	22214	
unique	22214	7758	17754	10854	3	7	13	

```

top      99-9963543    Cassie    Ducker   5/9/1972    Male    White  Engineering
freq                 1          12         7          9    11288     6328       6686

                jobtitle    location hire_date \
count            22267      22214    22214
unique           185        2        7016
top    Research Assistant II Headquarters  3/4/2003
freq              754      16715       10

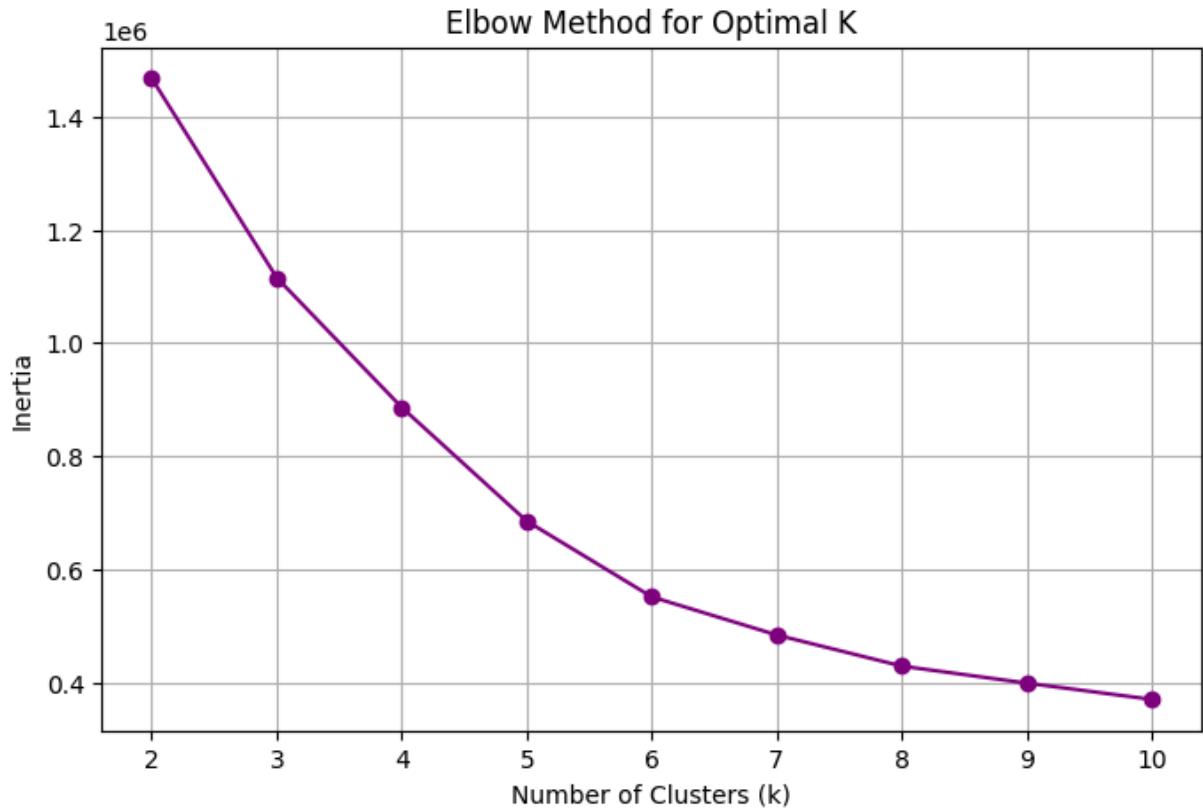
                termdate location_city location_state
count            3929      22214    22214
unique           3929        77        7
top    2012-12-10 14:29:59 UTC      Cleveland        Ohio
freq                  1      16871    18025
-----
Dataset Shape
(37408, 13)
-----
```

```

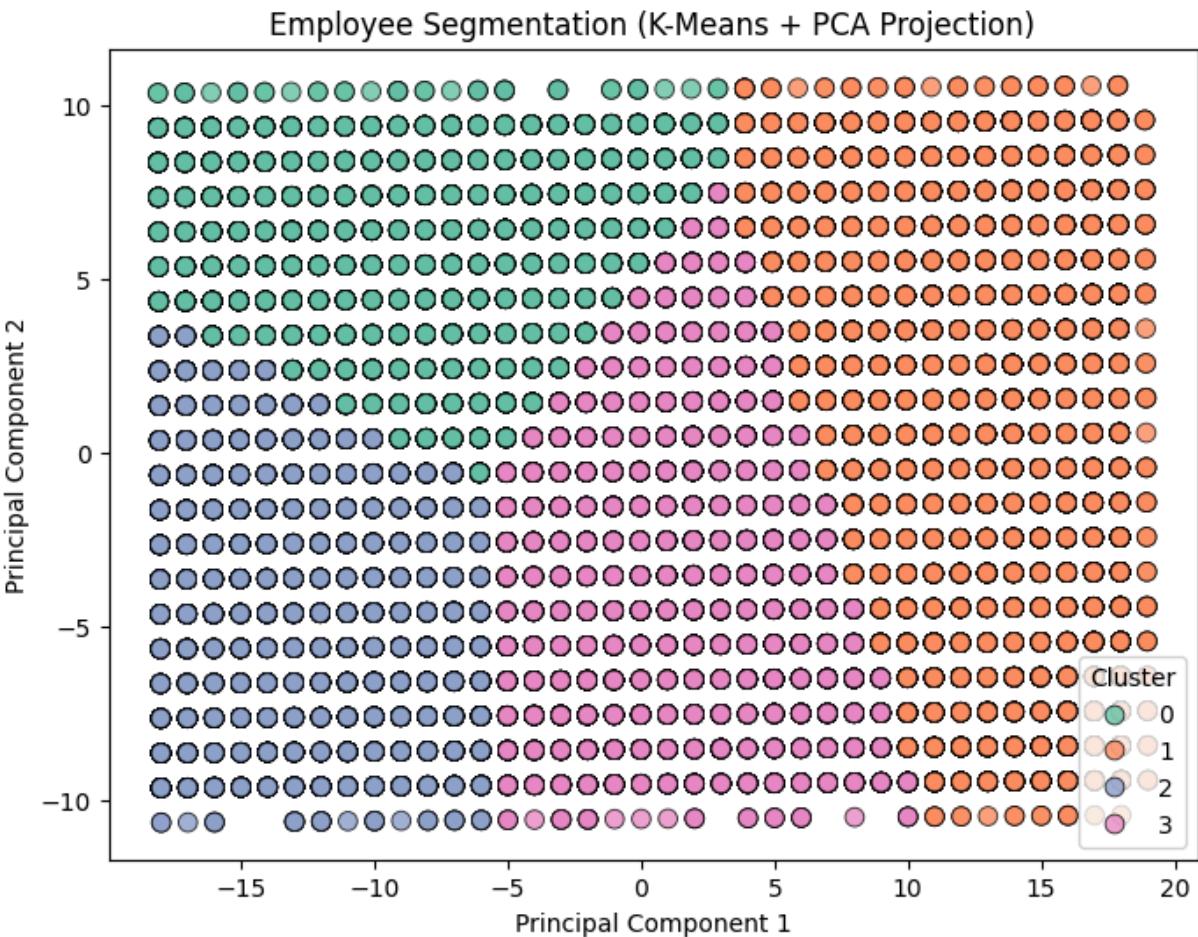
1 from sklearn import cluster
2 from sklearn.preprocessing import StandardScaler
3
4
5 # # EDA
6
7 # print('Check Missing Values')
8 # print(hr_df.isnull().sum())
9 # print("-----")
10
11 # View & Vizualize Missing Cols
12
13 # print('View & Vizualize Missing Cols')
14 # missing_perct_cols = (hr_df.isnull().sum()/hr_df.shape[0])*100
15 # print(missing_perct_cols)
16
17 # Viz - Using Heatmap to indicate the percentage of cols missing in my dat
18
19 # plt.figure(figsize=(10, 6))
20 # sns.heatmap(hr_df.isnull(), cbar=False, cmap='viridis')
21 # plt.title('Missing Values Heatmap')
22 # plt.show()
23
24 # Drop & Fill Missing Cols
25
26 # print('Drop Rows that are completely Empty')
27 # hr_df.dropna(how='all', inplace=True)
28 # print(hr_df)
29 # print("-----")
30
31 # hr_df.drop(['id', 'first_name','last_name'], axis=1, inplace=True, error
32
33 # # Fill Missing Cols - Re-evaluating this section to ensure date columns
```

```
34
35 # print('Fill Missing Cols')
36 # hr_df['department'] = hr_df['department'].fillna('Unknown')
37 # print(hr_df)
38 # print("-----")
39
40 # print(hr_df.columns.tolist()) #to check cols
41 # # print(hr_df['termdate'].unique) # to verify that the fill missing cols
42
43 # Convert Date Cols to Datetime
44
45 # hr_df['birthdate'] = pd.to_datetime(hr_df['birthdate'], errors='coerce')
46 # hr_df['hire_date'] = pd.to_datetime(hr_df['hire_date'], errors='coerce')
47 # hr_df['termdate'] = pd.to_datetime(hr_df['termdate'], errors='coerce') #
48
49 # Handle Missing Values in Date Columns AFTER Conversion
50
51 # For numerical calculations like age and years of service, fill with medi
52
53 # median_birth = hr_df['birthdate'].median() # Replace missing birthdate &
54 # median_hire = hr_df['hire_date'].median()
55
56 # hr_df['birthdate'].fillna(median_birth, inplace=True)
57 # hr_df['hire_date'].fillna(median_hire, inplace=True)
58
59 # # Sanity Check
60 # print("Number of missing birthdates after fill:", hr_df['birthdate'].isr
61 # print("Number of missing hire_dates after fill:", hr_df['hire_date'].isr
62
63
64 ##For termdate, where 'Active' was used, NaT represents 'Active' after coe
65
66 # Feature Engineering (to creates additional cols from the dataset for sig
67
68 # today = pd.to_datetime('today').replace(tzinfo=None)
69 # hr_df['age'] = (today - hr_df['birthdate']).dt.days // 365
70 # # Assuming 'hire_date' will now be in datetime format after dropping NaT
71 # hr_df['years_of_service'] = (today - hr_df['hire_date']).dt.days // 365
72
73 ## Adding Safety Check to ensure no NaT remain in my numerical syntax:
74
75 # hr_df['age'] = hr_df['age'].fillna(hr_df['age'].median())
76 # hr_df['years_of_service'] = hr_df['years_of_service'].fillna(hr_df['year
77
78 # Create new employment status column based on termdate NaT values
79
80 # hr_df['employment_status'] = hr_df['termdate'].apply(lambda x: 'Active'
81
82
83 # Drop any date_time cols:
84
```

```
85 # # Drop datetime columns that were used to calculate numeric features
86 # hr_df.drop(columns=['birthdate', 'hire_date', 'termdate'], inplace=True,
87
88 # Verify no datetime columns remain in my dataset:
89
90 # print("Remaining datetime columns:", hr_df.select_dtypes(include=[ 'datet
91
92 # Encode and Scale the Dataset:
93
94 # -- Select the Categorical Dataset ---
95
96 categorical_cols = ['gender', 'race', 'department', 'jobtitle', 'location_city
97 numerical_cols = ['age', 'years_of_service']
98
99 # -- One-Hot Encode Categorical Cols ---
100
101 # Applying one-hot encoding to the specified categorical columns
102 hot_codes = pd.get_dummies(hr_df, columns=categorical_cols, drop_first=True)
103
104
105 # --- Scale Numerical Features---using Scaler function from Sklearn
106 # scaler = StandardScaler()
107 # hot_codes[numerical_cols] = scaler.fit_transform(hot_codes[numerical_col
108
109 # print("Conversion of Hot Encoding & Scale Codes:", hot_codes.shape)
110 # print(hot_codes.head())
111
112 # Conduct Clustering on my dataset:
113
114 inertia = []
115 Cluster = range(2,11) # Use range to iterate from 2 to 10 clusters
116
117 for cluster in Cluster:
118     k_means = KMeans(n_clusters=cluster, random_state=42, n_init=10) # Add
119     k_means.fit(hot_codes)
120     inertia.append(k_means.inertia_)
121
122
123 # Plot Elbow Curve
124 plt.figure(figsize=(8,5))
125 plt.plot(Cluster, inertia, 'o-', color='purple')
126 plt.title('Elbow Method for Optimal K')
127 plt.xlabel('Number of Clusters (k)')
128 plt.ylabel('Inertia')
129 plt.grid(True)
130 plt.show()
```



```
1 from sklearn.decomposition import PCA
2
3 # --- PCA Transformation ---
4 pca = PCA(n_components=2)
5 pca_result = pca.fit_transform(hr_df)
6
7 # Create a DataFrame for PCA components and clusters
8 pca_df = pd.DataFrame(pca_result, columns=['PCA1', 'PCA2'])
9 pca_df['Cluster'] = hr_df['KMeans_Cluster']
10
11 # --- Improved PCA Visualization ---
12 plt.figure(figsize=(8,6))
13 sns.scatterplot(
14     x='PCA1', y='PCA2',
15     hue='Cluster', data=pca_df,
16     palette='Set2', s=60, alpha=0.8, edgecolor='k'
17 )
18 plt.title("Employee Segmentation (K-Means + PCA Projection)")
19 plt.xlabel("Principal Component 1")
20 plt.ylabel("Principal Component 2")
21 plt.legend(title='Cluster', loc='best')
22 plt.show()
23
```



Results Interpretation for the above Cluserings / Graphs (K-MEANS & PCA)

The Elbow Method graph showed a clear inflection point at $k = 4$, suggesting that four clusters provide the most meaningful segmentation of the HR dataset. Beyond four clusters, the reduction in inertia became marginal, confirming that additional clusters would not significantly improve model performance. Thus, $k = 4$ was selected as the optimal number of employee groups for analysis.

After applying K-Means clustering with $k = 4$ and visualizing through PCA, distinct employee clusters emerged. The PCA projection revealed well-separated regions, indicating that the selected features—such as age, years of service, department, job title, and employment status—contribute effectively to distinguishing employee profiles. The clusters show clear differentiation, with each group representing unique employee characteristics and work patterns.

In practical terms, these clusters may reflect different workforce segments—for example, long-tenured employees nearing retirement, mid-career staff with stable performance, new hires or early-career employees, and remote or specialized roles. HR managers could leverage these insights to tailor engagement and retention strategies to each group, improve onboarding for new hires, and design development programs that address the needs of mid-career employees. Overall, the clustering results provide a data-driven foundation for more targeted and equitable workforce planning.

```
1 from sklearn.metrics import silhouette_score
2
3 # Compute Silhouette Score for the final model
4 silhouette_avg = silhouette_score(hr_df['KMeans_Cluster'], hr_df)
5 print(f"Silhouette Score: {silhouette_avg:.3f}")
6
7
8 # Evaluate multiple cluster counts for comparison
9
10 for k in range(2, 7):
11     kmeans = KMeans(n_clusters=k, random_state=42)
12     labels = kmeans.fit_predict(hr_df)
13     score = silhouette_score(hr_df, labels)
14     print(f"k={k}: Silhouette Score = {score:.3f}")
15
```

```
Silhouette Score: 0.322
k=2: Silhouette Score = 0.456
k=3: Silhouette Score = 0.373
k=4: Silhouette Score = 0.322
k=5: Silhouette Score = 0.333
k=6: Silhouette Score = 0.341
```

At first, when I ran the silhouette score and got 0.3222. I needed to compare the score using other clusters based on a range, which I was aiming for a score closer to +1. However, the additional Silhouette Scores indicate how well the data points fit within their assigned clusters for different numbers of clusters (k)

The scores are reflective of how distinctly employees can be grouped based on their demographic and job-related features. In real-life terms, the higher score at k=2 (0.456) may represent two broad employee categories — for example, new hires versus long-tenured staff.

As the number of clusters increases (k=3–6), scores slightly decline, meaning the model starts identifying smaller subgroups such as mid-career employees, remote workers, or specialized job roles, which naturally overlap in characteristics. Choosing k=4 (0.322) captures this complexity realistically — mirroring how organizations often have multiple

overlapping workforce segments rather than completely distinct groups. These clusters can be used by HR teams to design targeted initiatives such as onboarding programs for new hires, career development plans for mid-level employees, or retention strategies for experienced staff.