

```
1 # Import the various libraries:  
2  
3 from tensorflow.keras.datasets import cifar10  
4 import numpy as np  
5  
6 # Load the Dataset from the in-bult library  
7  
8 (X_train, y_train), (X_test, y_test) = cifar10.load_data() # this syntax fe  
9  
10 # Print the Dataset  
11  
12 print(X_train.shape)  
13 print(y_train.shape)  
14 print(X_test.shape)  
15 print(y_test.shape)  
16
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 ━━━━━━ 3s 0us/step
(50000, 32, 32, 3)
(50000, 1)
(10000, 32, 32, 3)
(10000, 1)

```
1 # Normalize the Dataset (This is part of the preprocessing stage the data is  
2  
3 X_train_norm = X_train.astype('float32')/255.0  
4 X_test_norm = X_test.astype('float32') / 255.0  
5  
6 print("The Normalize Dataset Shape for:", X_train_norm.shape) # this syntax  
7 print("The Normalize Dataset Shape for:", X_test_norm.shape)  
8  
9 # Sample dataset for training and testing. Using the full dataset(50,000) c  
10  
11 training_values = 5000 # this syntax reflects the proportion of the dataset  
12 testing_values = 2000  
13  
14  
15 # Reduce Images  
16  
17 X_train_small = X_train_norm[:training_values]  
18 X_test_small = X_test_norm[:testing_values]  
19  
20  
21 # Reduce labels  
22  
23 y_train_small = y_train[:training_values] # the dataset in the label must a  
24 y_test_small = y_test[:testing_values]  
25
```

```
1 # Import the required libraries
2 import numpy as np
3 import pandas as pd
4 import tensorflow as tf
5 import tensorflow_hub as hub
6 import time
7 import os
8 import cv2
9 import random
10 from PIL import Image
11
12 # Set the path to the dataset
13 dataset_path = "C:/Users/Asus/Downloads/flower_recognition"
14
15 # Load the dataset
16 df = pd.read_csv("C:/Users/Asus/Downloads/flower_recognition/flowers.csv")
17
18 # Preprocess the dataset
19 # 1. Convert the image paths to absolute paths
20 # 2. Extract the flower names from the file names
21 # 3. Create a mapping between flower names and indices
22
23 flower_names = df['name'].unique()
24 flower_index_map = {name: index for index, name in enumerate(flower_names)}
25
26 # Flatten the Images for Random Forest - this ensures that the images in the
27 # dataset are represented as 1D arrays
28 X_train_flat = X_train_small.reshape(training_values, -1)
29 X_test_flat = X_test_small.reshape(testing_values, -1)
30
31 print(X_train_flat.shape, y_train_small)
32 print(X_test_flat.shape, y_test_small.shape)
```

```
The Normalize Dataset Shape for: (50000, 32, 32, 3)
The Normalize Dataset Shape for: (10000, 32, 32, 3)
(5000, 3072) [[6]
[9]
[9]
...
[5]
[4]
[6]]
(2000, 3072) (2000, 1)
```

```
1 # The Model ( Train the Model to ascertain if it performance on the dataset
2
3 # Import the Random Forest Model
4
5 from sklearn.ensemble import RandomForestClassifier
6
7 # this is the model used for training
8 random_forest_model = RandomForestClassifier(n_estimators=100, random_state=
9
10 # train the model
11 random_forest_model.fit(X_train_flat, y_train_small.ravel())
12
13 print("Training Complete")
```

Training Complete

```
1 # Predict the Model
2
3 y_pred = random_forest_model.predict(X_test_flat)
4
5 # Evaluate the Model
6
7 from sklearn.metrics import accuracy_score, classification_report
8
9 accuracy = accuracy_score(y_test_small, y_pred)
10
11 print("Random Forest Accuracy Score:", accuracy)
12 print("Classification Report:\n", classification_report(y_test_small, y_pred))
13
14 import seaborn as sns
15 import matplotlib.pyplot as plt
16 from sklearn.metrics import confusion_matrix
```

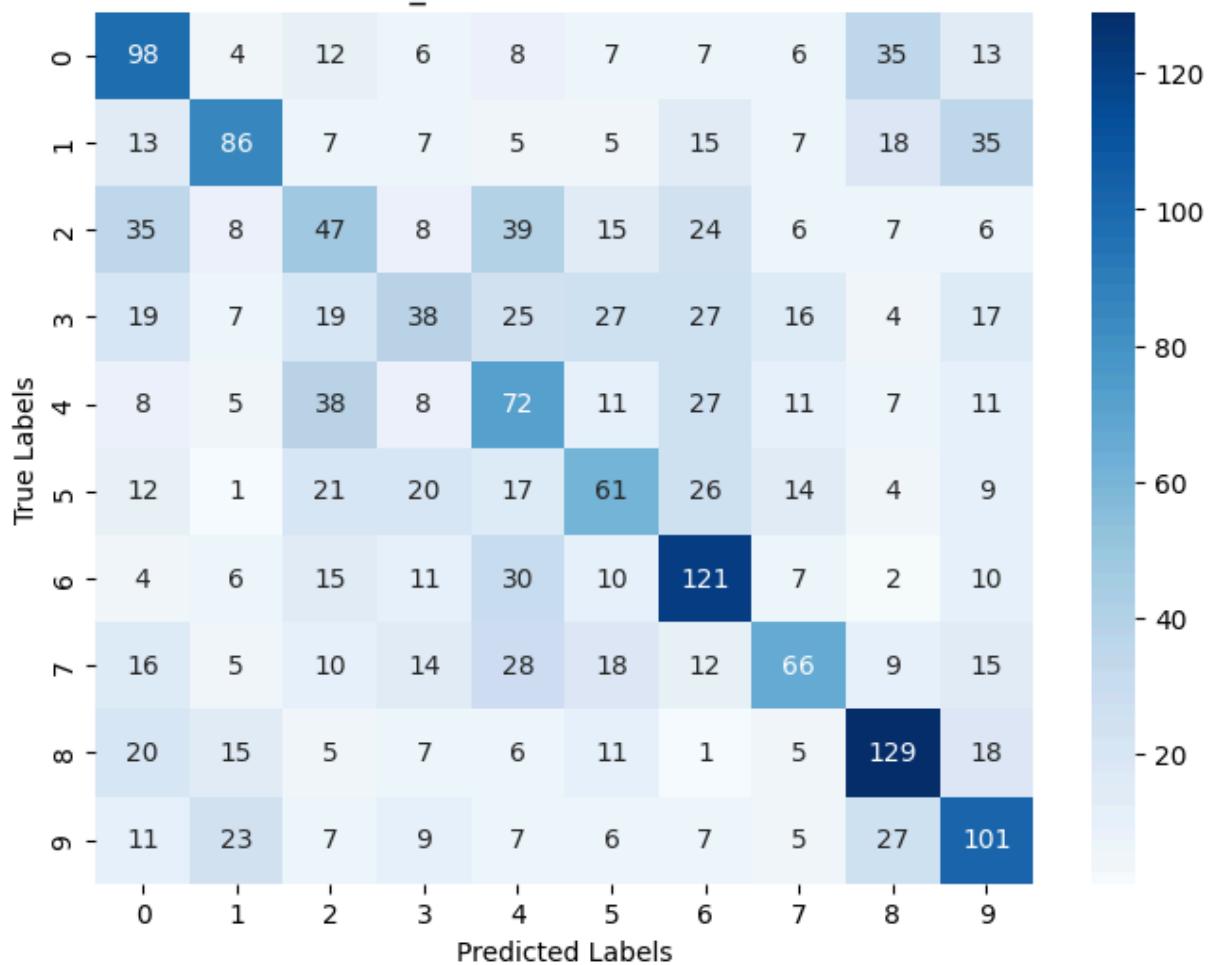
```
17
18 confix_matrix = confusion_matrix(y_test_small, y_pred)
19
20 plt.figure(figsize=(8,6))
21 sns.heatmap(confix_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=r
22 plt.xlabel("Predicted Labels")
23 plt.ylabel("True Labels")
24 plt.title("Confusion Matrix _ Random Forest on CIFAR10 Dataset")
25 plt.show()
```

Random Forest Accuracy Score: 0.4095

Classification Report:

	precision	recall	f1-score	support
0	0.42	0.50	0.45	196
1	0.54	0.43	0.48	198
2	0.26	0.24	0.25	195
3	0.30	0.19	0.23	199
4	0.30	0.36	0.33	198
5	0.36	0.33	0.34	185
6	0.45	0.56	0.50	216
7	0.46	0.34	0.39	193
8	0.53	0.59	0.56	217
9	0.43	0.50	0.46	203
accuracy			0.41	2000
macro avg	0.40	0.41	0.40	2000
weighted avg	0.41	0.41	0.40	2000

Confusion Matrix _ Random Forest on CIFAR10 Dataset



```

1 from sklearn.model_selection import GridSearchCV
2
3 #Fining the Parameters

```

```

4
5 param_grid = {
6     'n_estimators': [50, 100, 150],           # I used a small range of trees to
7     'max_depth': [None, 10, 20, 30],         # shallow vs deeper
8     'min_samples_split': [2, 5, 10],        # small options
9     'min_samples_leaf': [1]                  # fixed to reduce combination
10 }
11
12 # Set - up the Base
13 random_forest_model = RandomForestClassifier(random_state=42, n_jobs=-1)
14
15 grid_search = GridSearchCV(                 # this is integration of the GridSearchCV
16     estimator=random_forest_model,
17     param_grid=param_grid,
18     cv=3,                                # 3 fold cross validation
19     n_jobs=-1,                            # use all the cpu cores
20     verbose=2)                           # show progress
21
22 grid_search.fit(X_train_flat, y_train_small.ravel())
23
24 #Print the best GridSearch Result:
25
26 print("Best Parameters:", grid_search.best_params_)
27 print("Best Score:", grid_search.best_score_)

```

Fitting 3 folds for each of 36 candidates, totalling 108 fits
 Best Parameters: {'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 1}
 Best Score: 0.40140303271878636

```

1 # Train the Random Forest Model based on the GridSearchCV Results.
2 ## GridSearchCV does not train my model, instead it observed the model
3 ## based on the parameters that were established and provided me with the best
4
5 # Retrieve the best parameters from the GridSearchCV
6
7 best_params = grid_search.best_params_
8 print("Best Parameters:", best_params)
9
10 # Train the final model
11
12 final_model = RandomForestClassifier(
13     n_estimators = best_params['n_estimators'],
14     max_depth= best_params['max_depth'],
15     min_samples_split = best_params['min_samples_split'],
16     min_samples_leaf=best_params['min_samples_leaf'],
17     random_state =42,
18     n_jobs = -1,
19     verbose = 0 # the output from my RF indicated the logs of the training
20 )
21

```

```

22
23 # # Fit the final model on the full reduce text/dataset
24
25 final_model.fit(X_train_flat, y_train_small.ravel())
26
27 print("Final Model Training Complete")
28

```

Best Parameters: {'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 1}
Final Model Training Complete

```

1 # Evaluate the Model
2
3 y_pred_final = final_model.predict(X_test_flat)
4
5 accuracy = accuracy_score(y_test_small, y_pred_final)
6
7 print("Random Forest Accuracy Score:", accuracy)
8 print("Classification Report:\n", classification_report(y_test_small, y_pre

```

Random Forest Accuracy Score: 0.403

Classification Report:

	precision	recall	f1-score	support
0	0.39	0.46	0.42	196
1	0.56	0.41	0.48	198
2	0.24	0.21	0.22	195
3	0.35	0.18	0.24	199
4	0.30	0.34	0.32	198
5	0.30	0.31	0.31	185
6	0.44	0.53	0.48	216
7	0.45	0.35	0.39	193
8	0.53	0.62	0.57	217
9	0.44	0.56	0.49	203
accuracy			0.40	2000
macro avg	0.40	0.40	0.39	2000
weighted avg	0.40	0.40	0.40	2000

This test is based on the above evaluation that was ran for the final RF Model

▼ Results & Model Interpretation:

The final tuned Random Forest model achieved an accuracy of approximately 40.3% on the CIFAR-10 test set. This level of performance is expected for a classical machine learning algorithm applied to image data. CIFAR-10 contains complex, high-dimensional images (32×32×3 pixels), and Random Forests treat each pixel as an independent feature

without understanding spatial patterns or textures. As a result, the model is limited in its ability to capture visual structure compared to convolutional neural networks (CNNs).

Across individual classes, the model performed better on objects with more distinct shapes and colors, such as automobiles, trucks, ships, and frogs. These classes achieved relatively higher precision and recall. In contrast, classes such as cats, dogs, deer, and birds had lower performance because they share similar shapes and color patterns, making them harder for a pixel-based model to separate.

Overall, the results demonstrate that while Random Forests can provide some classification ability, they are not well-suited for complex image tasks. The experiment highlights the importance of using specialized deep learning architectures for image recognition. Nevertheless, the model behaves consistently and validates the effectiveness of preprocessing, parameter tuning, and evaluation steps in the machine learning workflow.

```
1
2 from PIL import Image
3
4 # Prediction on new image
5
6 cifar10_classes = [
7     "airplane", "automobile", "bird", "cat", "deer",
8     "dog", "frog", "horse", "ship", "truck"
9 ]
10
11 def predict_new_image(image_path, model):
12     img = Image.open(image_path).convert('RGB') # load the image and conver
13     img = img.resize((32, 32)) # this syntax resize the image as per the re
14     img_array = np.array(img) # convert the image to an array
15     img_array = img_array.astype('float32') / 255.0 # normalize the image
16     img_flat = img_array.reshape(1, -1) # flatten the image leads to equal
17     predicted_model = model.predict(img_flat) # predict the image
18     predicted_class = cifar10_classes[predicted_model[0]] # convert predict
19     return predicted_class
```

```
1 from google.colab import files # this syntax allows me to upload a picture
2 uploaded = files.upload()
3
4 filename = list(uploaded.keys())[0]
5 image_path = filename
6
7 result = predict_new_image(image_path, final_model)
8 print("Predicted Class:", result)
9
```

Choose Files s4vstarget.png

s4vstarget.png(image/png) - 183335 bytes, last modified: 8/4/2025 - 100% done

Saving s4vstarget.png to s4vstarget.png

Predicted Class: airplane

```
1 from sklearn.svm import SVC # this library is call the Support Vector Machi
2
3 svm_model = SVC(kernel='linear')
4
5 print("Training Model SVC")
6 svm_model.fit(X_train_flat, y_train_small.ravel())
7 print("Training Complete")
```

Training Model SVC

Training Complete

```
1 # Predict the Model
2
3 y_pred_svm = svm_model.predict(X_test_flat)
4
5 # Print the Reports - Accuracy, SVM & Classification Reports
6
7 svm_model_accuracy = accuracy_score(y_test_small, y_pred_svm)
8
9 print("SVM Accuracy Score:", svm_model_accuracy)
10 print("Classification Report:\n", classification_report(y_test_small, y_p
11
12 print("Random Forest Accuracy Score:", accuracy)
13 print("Classification Report:\n", classification_report(y_test_small, y_p
14
```

SVM Accuracy Score: 0.3

Classification Report:

	precision	recall	f1-score	support
0	0.27	0.39	0.32	196
1	0.39	0.38	0.39	198
2	0.24	0.31	0.27	195
3	0.20	0.21	0.20	199
4	0.25	0.26	0.25	198
5	0.16	0.15	0.15	185
6	0.36	0.27	0.31	216
7	0.32	0.27	0.29	193
8	0.45	0.46	0.45	217
9	0.40	0.28	0.33	203
accuracy			0.30	2000
macro avg	0.30	0.30	0.30	2000
weighted avg	0.31	0.30	0.30	2000

Random Forest Accuracy Score: 0.403

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0	0.27	0.39	0.32	196
1	0.39	0.38	0.39	198
2	0.24	0.31	0.27	195
3	0.20	0.21	0.20	199
4	0.25	0.26	0.25	198
5	0.16	0.15	0.15	185
6	0.36	0.27	0.31	216
7	0.32	0.27	0.29	193
8	0.45	0.46	0.45	217
9	0.40	0.28	0.33	203
accuracy			0.30	2000
macro avg		0.30	0.30	2000
weighted avg		0.31	0.30	2000

Deployment Strategy:

Deploying an image classification model involves preparing the trained system so it can be used reliably in real-world applications. Although this project uses a classical machine learning approach, the deployment workflow follows the same principles used in modern ML systems. The strategy below outlines how the Random Forest model can be operationalized in a production environment.

1. Model Packaging

After training, the final Random Forest model should be saved using a serialization format such as pickle or joblib. This allows the model to be reloaded without retraining. The preprocessing steps—image resizing, normalization, and flattening—must also be packaged with the model so the same transformations are applied consistently during inference.

2. Building an Inference Pipeline

A lightweight prediction pipeline is needed to accept new images and return output: Receive the uploaded image (e.g., JPG/PNG). Convert it to RGB, resize it to 32×32, and normalize pixel values. Flatten the image into a 1D array of 3,072 features. Pass the processed array to the saved model for prediction. Map the numeric prediction to the corresponding CIFAR-10 class label. This ensures the model handles new images in the exact same manner as during training.

3. Deployment Options

There are multiple ways the model could be deployed depending on system requirements:

- Web API Deployment : The model can be hosted behind a simple REST API using tools such as:
 - Flask
 - FastAPI
 - Django REST Framework
 Clients can send images to the API endpoint, and the server returns the predicted class.
- Cloud Deployment: The

model can be deployed using cloud services such as:

- AWS Lambda
- Google Cloud Functions
- Azure App Service

These services are automatically scaled based on user demand and simplify maintenance.

Local or Desktop Application:

For offline use, the model can be integrated into a:

- desktop application
- mobile app
- command-line tool