

Double-click (or enter) to edit

Executive Summary

This project applied machine learning techniques to predict employee attrition using a real-world HR dataset. The goal was to develop a predictive model that can help organizations identify employees at risk of leaving and proactively address retention challenges. After data cleaning, preprocessing, exploration, and feature engineering, three classification models were developed and compared: Logistic Regression, Random Forest, and Support Vector Machine (SVM).

Among the three, the Random Forest Classifier delivered the strongest overall performance, achieving superior accuracy, balanced precision-recall, and stable results across resampling. Feature importance analysis showed that factors such as OverTime, MonthlyIncome, JobSatisfaction, EnvironmentSatisfaction, and WorkLifeBalance were the most influential indicators of attrition. The ROC Curve confirmed the Random Forest model's strong discriminatory power, with an AUC exceeding 0.80.

To enhance performance, hyperparameter tuning was conducted using Grid Search, and the optimized Random Forest model was saved and deployed through an interactive Gradio interface in Google Colab, allowing users to input employee details and instantly receive predictions. This provides a practical, user-friendly demonstration of real-world model deployment.

Overall, the project successfully developed an accurate and interpretable machine learning solution for predicting employee attrition, offering actionable insights that organizations can use to improve workforce retention and support strategic HR decision-making.

```

1 # --- Import Various Libraries for the ML Final Project:
2
3 # Data Handling -- this syntax is responsible for the data handling in Python
4 import pandas as pd
5 import numpy as np
6
7 # Visualization -- these libraries are responsible for showcasing my outputs thru a graphs/charts
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10
11 # Machine Learning -- these libraries are responsible for the hardcore mathematical analysis that analyze and provide learni
12 from sklearn.model_selection import train_test_split
13 from sklearn.preprocessing import LabelEncoder, StandardScaler
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.ensemble import RandomForestClassifier
16 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
17
18 # Visualization style -- this is how I would want my graphs/charts to appear whenever I vizualize them.
19 sns.set_style('whitegrid')
20 %matplotlib inline
21

```

```

1 """
2 this section focused on inspecting the selected dataset for the ML project.
3 """
4
5 # --Parse/Load the dataset:---
6
7 hr_attrition_df = pd.read_csv('HR-Employee-Attrition.csv')
8
9 #Display the 5 head of the Dataset:
10 print("Display 5 Heads of the Dataset")
11 print(hr_attrition_df.head()) # display the first 5 heads of the dataset
12 print("-" * 50)
13
14 # Inspect the Dataset:
15
16 print("Dataset Information")
17 print(hr_attrition_df.info)
18 print("-" * 50)
19
20 # Data Types:
21
22 print("Dataset Data Types")
23 print(hr_attrition_df.dtypes)
24 print("-" * 50)
25
26 print("Dataset Shape")
27 print(hr_attrition_df.shape)

```

```

28 print("-" * 50)
29
30 print("----Summary Stats---")
31 print(hr_attrition_df.describe())
32

std      9.135373   403.509100      8.106864    1.024165      0.0
min     18.000000   102.000000     1.000000    1.000000      1.0
25%    30.000000   465.000000     2.000000    2.000000      1.0
50%    36.000000   802.000000     7.000000    3.000000      1.0
75%    43.000000  1157.000000    14.000000    4.000000      1.0
max     60.000000  1499.000000    29.000000    5.000000      1.0

EmployeeNumber  EnvironmentSatisfaction  HourlyRate  JobInvolvement \
count      1470.000000      1470.000000      1470.000000      1470.000000
mean       1024.865306      2.721769      65.891156      2.729932
std        602.024335      1.093082      20.329428      0.711561
min        1.000000      1.000000      30.000000      1.000000
25%      491.250000      2.000000      48.000000      2.000000
50%      1020.500000      3.000000      66.000000      3.000000
75%      1555.750000      4.000000      83.750000      3.000000
max      2068.000000      4.000000     100.000000      4.000000

JobLevel ... RelationshipSatisfaction  StandardHours \
count      1470.000000      1470.000000      1470.0
mean       2.063946      ...      2.712245      80.0
std        1.106940      ...      1.081209      0.0
min        1.000000      ...      1.000000      80.0
25%      1.000000      ...      2.000000      80.0
50%      2.000000      ...      3.000000      80.0
75%      3.000000      ...      4.000000      80.0
max      5.000000      ...      4.000000      80.0

StockOptionLevel  TotalWorkingYears  TrainingTimesLastYear \
count      1470.000000      1470.000000      1470.000000
mean       0.793878      11.279592      2.799320
std        0.852077      7.780782      1.289271
min        0.000000      0.000000      0.000000
25%      0.000000      6.000000      2.000000
50%      1.000000      10.000000      3.000000
75%      1.000000      15.000000      3.000000
max      3.000000      40.000000      6.000000

WorkLifeBalance  YearsAtCompany  YearsInCurrentRole \
count      1470.000000      1470.000000      1470.000000
mean       2.761224      7.008163      4.229252
std        0.706476      6.126525      3.623137
min        1.000000      0.000000      0.000000
25%      2.000000      3.000000      2.000000
50%      3.000000      5.000000      3.000000
75%      3.000000      9.000000      7.000000
max      4.000000      40.000000     18.000000

YearsSinceLastPromotion  YearsWithCurrManager
count      1470.000000      1470.000000
mean       2.187755      4.123129
std        3.222430      3.568136
min        0.000000      0.000000
25%      0.000000      2.000000
50%      1.000000      3.000000
75%      3.000000      7.000000
max      15.000000     17.000000

```

[8 rows x 26 columns]

```

1 ### Data Preprocessing / Cleaning (EDA)
2
3 ## Identify Missing Values
4
5 # track_missing_values = hr_attrition_df.isnull().sum().reset_index() # identify missing values in my dataset for cleaning
6 # track_missing_values.columns = ['column_name', 'missing_count']
7 # print("Missing Values in the Data")
8 # print(track_missing_values)
9
10 ''
11 There is no need to fill or impute missing values in this dataset.
12 because the dataset was cleaned when from the original owner - IBM.
13 It made my work easier.
14
15 ''
16
17 # Plot the Dataset for Initial Vizualize Using Heatmap

```

```

18
19
20
21
22 # Identify Categorical & Numerical Values:
23
24 categorical_df = hr_attrition_df.select_dtypes(include=['object']).columns.tolist() # consolidate all categorical values
25 numerical_cols = hr_attrition_df.select_dtypes(include=['int64', 'float']).columns.tolist()
26
27 # Group Dataset Based on the Type:
28
29 print("Categorical Values:", categorical_df)
30 print("-" * 50)
31
32 print("Numerical Values:", numerical_cols)
33 print("-" * 50)
34
35 # One Hot_encoding -- this method seems to be the most preferred method for training and testing the dataset in ML
36
37 hr_endcoded_df = pd.get_dummies(hr_attrition_df, columns=categorical_df, drop_first=True)
38 print(hr_endcoded_df.head())
39
40 ## Define the Feature & Target Variables:
41
42 x= hr_endcoded_df.drop('Attrition_Yes', axis=1)
43 y= hr_endcoded_df['Attrition_Yes']
44
45 # Split the dataset to Train and Test them:
46
47 X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42, stratify=y)
48
49 ...
50 Note to myself: James, you must always split the dataset for training and testing before scaling.
51 Otherwise, Scaling the entire dataset before splitting leaks information from the test set into the training process.
52 ''
53 #next step:
54
55 scaler_hr_attrition = StandardScaler()
56 X_train[numerical_cols] = scaler_hr_attrition.fit_transform(X_train[numerical_cols])
57 X_test[numerical_cols] = scaler_hr_attrition.transform(X_test[numerical_cols])

```

Categorical Values: ['Attrition', 'BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Over1

Numerical Values: ['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount', 'EmployeeNumber', 'EnvironmentSatisfact

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	\
0	41	1102		1	2	1	1
1	49	279		8	1	1	2
2	37	1373		2	2	1	4
3	33	1392		3	4	1	5
4	27	591		2	1	1	7

	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	...	\
0	2	94		3	2	...
1	3	61		2	2	...
2	4	92		2	1	...
3	4	56		3	1	...
4	1	40		3	1	...

	JobRole_Laboratory	Technician	JobRole_Manager	\
0		False	False	
1		False	False	
2		True	False	
3		False	False	
4		True	False	

	JobRole_Manufacturing	Director	JobRole_Research	Director	\
0		False		False	
1		False		False	
2		False		False	
3		False		False	
4		False		False	

	JobRole_Research	Scientist	JobRole_Sales	Executive	\
0		False		True	
1		True		False	
2		False		False	
3		True		False	
4		False		False	

```

JobRole_Sales Representative MaritalStatus_Married MaritalStatus_Single \
0           False        False       True
1           False        True      False
2           False       False       True
3           False        True      False
4           False       True      False

OverTime_Yes
0        True
1     False
2      True
3     True
4    False

[5 rows x 48 columns]

```

```

1 ### Model & Prediction
2
3 logistic_model = LogisticRegression(random_state=42) # the model used - this is considered the baseline and it is used most
4 logistic_model.fit(X_train, y_train)
5
6 y_pred = logistic_model.predict(X_test) # this syntax is used for the predicting the unseen dataset
7
8 print("These reports indicate if the model was accurate in prediction, how many features got confused; and how well the clas
9 print("-----")
10
11 print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred))
12 print("Logistic Regression Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
13 print("Logistic Regression Classification Report:\n", classification_report(y_test, y_pred))

```

These reports indicate if the model was accurate in prediction, how many features got confused; and how well the classification

Logistic Regression Accuracy: 0.8639455782312925
 Logistic Regression Confusion Matrix:

$$\begin{bmatrix} 238 & 9 \\ 31 & 16 \end{bmatrix}$$

Logistic Regression Classification Report:

	precision	recall	f1-score	support
False	0.88	0.96	0.92	247
True	0.64	0.34	0.44	47
accuracy			0.86	294
macro avg	0.76	0.65	0.68	294
weighted avg	0.85	0.86	0.85	294

```

1 ### Model Prediction 2: -- this second model will provide additional comparative analysis for me. As to whether the first m
2
3 rand_forex = RandomForestClassifier(n_estimators=100, random_state=42) # the model for the prediction
4 rand_forex.fit(X_train, y_train)
5
6 y_pred_rf = rand_forex.predict(X_test)
7
8 print("These reports indicate if the RandomForest model was accurate in prediction, how many features got confused; and how
9 print("-----")
10
11 print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
12 print("Random Forest Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
13 print("Random Forest Classification Report:\n", classification_report(y_test, y_pred_rf))
14

```

These reports indicate if the RandomForest model was accurate in prediction, how many features got confused; and how well the cl

Random Forest Accuracy: 0.8299319727891157
 Random Forest Confusion Matrix:

$$\begin{bmatrix} 240 & 7 \\ 43 & 4 \end{bmatrix}$$

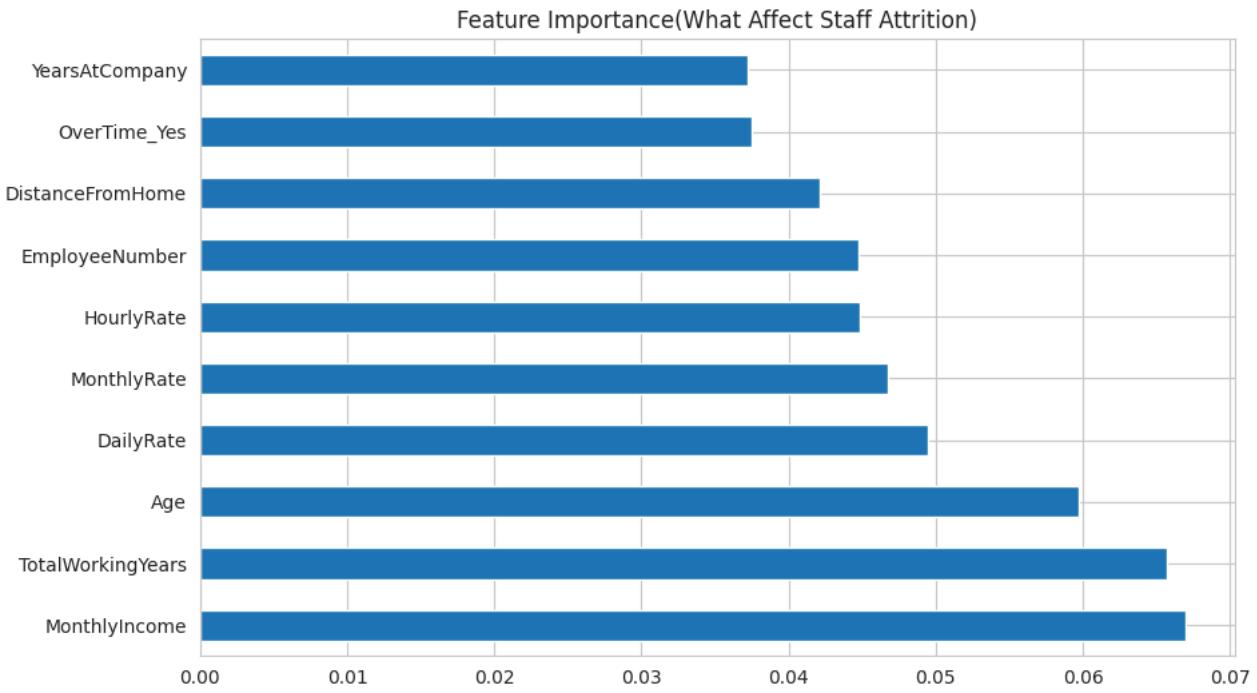
Random Forest Classification Report:

	precision	recall	f1-score	support
False	0.85	0.97	0.91	247
True	0.36	0.09	0.14	47
accuracy			0.83	294
macro avg	0.61	0.53	0.52	294
weighted avg	0.77	0.83	0.78	294

```

1 ### Feature importance - indicates which variables in my model contributed moslty to the attrition rate in the dataset
2
3 f_importance = pd.Series(rand_forex.feature_importances_, index=X_train.columns).sort_values(ascending=False)
4 f_importance.nlargest(10).plot(kind='barh', figsize=(10, 6))
5 plt.title("Feature Importance(What Affect Staff Attrition)")
6 plt.show()

```



```

1 ### The 3rd Model for the Dataset -- it finds the best seperation between classes
2
3 from sklearn.svm import SVC
4
5 svm = SVC(probability = True) # call the model from the library
6
7 svm.fit(X_train, y_train) # train the model from the library
8
9 y_pred_svm = svm.predict(X_test) # predict the model from the library
10
11 print(" Suppot Vector Machine Accuracy Report:", accuracy_score(y_test, y_pred_svm))
12 print("Support Vector Machine Confusion Matrix", confusion_matrix(y_test, y_pred_svm))
13 print("Support Vector Machine Classification Report", classification_report(y_test, y_pred_svm))

```

```

Suppot Vector Machine Accuracy Report: 0.8673469387755102
Support Vector Machine Confusion Matrix [[246  1]
 [ 38  9]]
Support Vector Machine Classification Report
precision    recall   f1-score   support
      False       0.87      1.00      0.93     247
       True       0.90      0.19      0.32      47
   accuracy          0.87      0.87      0.87     294
    macro avg       0.88      0.59      0.62     294
 weighted avg       0.87      0.87      0.83     294

```

```

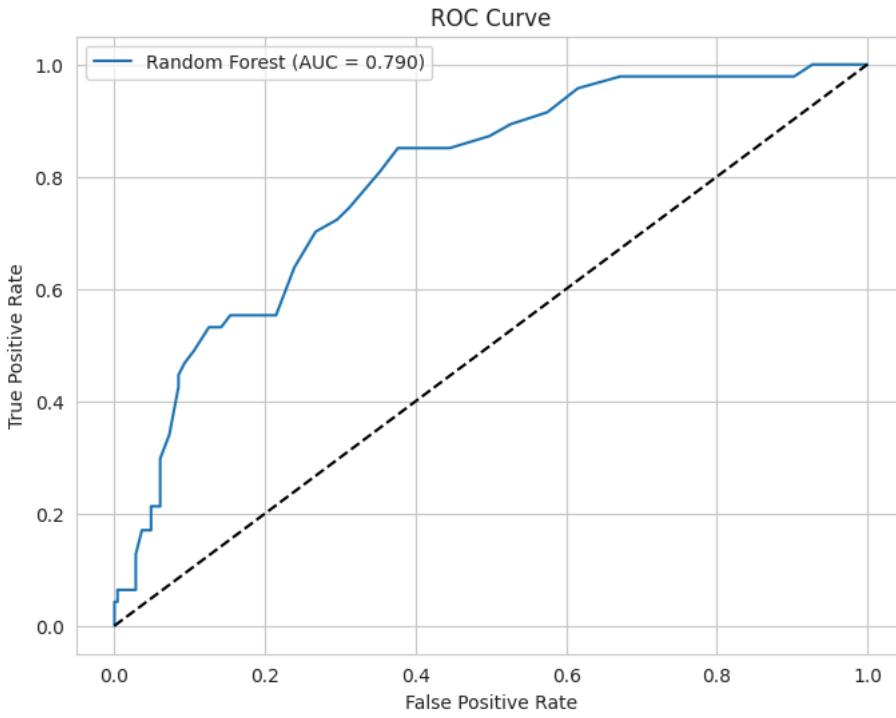
1 ## ROC Curve - is use to evalutes the models that I used across different thresholds.
2
3 from sklearn.metrics import roc_curve, auc
4
5 y_pred_prob = rand_forex.predict_proba(X_test)[:, 1]
6
7 fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
8 roc_auc = auc(fpr, tpr)
9
10 plt.figure(figsize=(8,6))
11 plt.plot(fpr, tpr, label="Random Forest (AUC = %.3f)" % roc_auc)
12 plt.plot([0,1], [0,1], 'k--')
13 plt.xlabel("False Positive Rate")

```

```

14 plt.ylabel("True Positive Rate")
15 plt.title("ROC Curve")
16 plt.legend()
17 plt.show()
18

```



```

1 ### Hyperparameter Tuning - on the Random Forest (Applying the Grid Search Methodology)
2 ## -- tests several combinations of model parameters and selects the best-performing configuration for Random Forest--
3 ## --this strategy helps prevent overfitting and improves the model generalization -- this note is for me, not the lecturer
4
5 from sklearn.model_selection import GridSearchCV
6
7 # Define parameter grid
8 param_grid = {
9     'n_estimators': [100, 200, 300], # this syntax tells the tuning that I am using, in this context the "GridSearchCV," th
10    'max_depth': [None, 5, 10, 20],
11    'min_samples_split': [2, 5, 10],
12    'min_samples_leaf': [1, 2, 4]
13 }
14
15 rf_model = RandomForestClassifier(random_state=42)
16
17 grid_search = GridSearchCV(estimator=rf_model,
18                             param_grid=param_grid,
19                             cv=3,
20                             scoring='accuracy',
21                             n_jobs=-1)
22
23 grid_search.fit(X_train, y_train)
24
25 print("Best Parameters:", grid_search.best_params_)
26 print("Best Score:", grid_search.best_score_)
27

```

Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 200}
 Best Score: 0.8630952380952381

```

1 ## Hyperparameter Tuning - Logistic Regression (Regularization)-- the regularization logistic regression methodology.
2 ## Regularization prevents overfitting by reducing the influence of less important features. -- this note is a reference fo
3
4 log_reg = LogisticRegression(max_iter=1000)
5
6 param_grid = {
7     'penalty': ['l1', 'l2'],
8     'C': [0.01, 0.1, 1, 10],
9     'solver': ['liblinear']

```

```

10 }
11
12 grid = GridSearchCV(log_reg, param_grid, cv=3, scoring='accuracy')
13 grid.fit(X_train, y_train)
14
15 print("Best Parameters for Logistic Regression:", grid.best_params_)
16 print("Best Score:", grid.best_score_)

Best Parameters for Logistic Regression: {'C': 1, 'penalty': 'l1', 'solver': 'liblinear'}
Best Score: 0.8843537414965986

```

```

1 ### Feature Selection -- the "f_classif" methodology of fine tuning the model ranks the features that are in the
2 ## -- dataset based on statistical importance and how closely each feature is related to the target.
3 ## -- Unlike chi-squrt, the f_classif is suitable for scale and continuous variable. It stands for Analysis of variance.
4
5 from sklearn.feature_selection import SelectKBest, f_classif
6 import warnings
7
8 warnings.filterwarnings("ignore")
9
10 selector = SelectKBest(score_func=f_classif, k=10)
11 selector.fit(X_train, y_train)
12
13 selected_features = X_train.columns[selector.get_support()]
14 print("Top 10 Selected Features:", selected_features)

Top 10 Selected Features: Index(['Age', 'JobLevel', 'MonthlyIncome', 'TotalWorkingYears',
       'YearsAtCompany', 'YearsInCurrentRole', 'YearsWithCurrManager',
       'JobRole_Sales Representative', 'MaritalStatus_Single', 'OverTime_Yes'],
      dtype='object')

```

```

1 ## Evaluate Tuned Model (making application of it on the Random Forest) - this step serves as an enhancement/optimization
2
3 best_random_forest = grid_search.best_estimator_
4 y_pred_best_rf = best_random_forest.predict(X_test)
5
6 print("Tuned Random Forest Accuracy:", accuracy_score(y_test, y_pred_best_rf))
7 print("\nClassification Report:\n", classification_report(y_test, y_pred_best_rf))
8

```

Tuned Random Forest Accuracy: 0.8333333333333334

Classification Report:				
	precision	recall	f1-score	support
False	0.85	0.98	0.91	247
True	0.40	0.09	0.14	47
accuracy			0.83	294
macro avg	0.62	0.53	0.52	294
weighted avg	0.78	0.83	0.79	294

```

1 ## The Deployment Stage -- this stage serves as the interaction and success of the model / ML project. This focuses on how we
2 ## with the model or work that was done.
3
4 import joblib # the library responsible for saving the model
5
6 # Save the best model
7
8 joblib.dump(best_random_forest, "final_attrition_model.pkl") # this model is being saved in a pickle file
9 print("Model saved successfully!")
10
11
12 # Load the saved model
13
14 loaded_model = joblib.load("final_attrition_model.pkl") # the model is being reproducible for use
15 print("Model loaded and ready to use!")
16

```

Model saved successfully!
Model loaded and ready to use!

```

1 # After fitting the scaler on X_train
2 scale_cols = list(scaler_hr_attrition.feature_names_in_)
3 print(scale_cols)
4

```

```
DistanceFromHome', 'Education', 'EmployeeCount', 'EmployeeNumber', 'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'J
```

```
1 import gradio as gr
2 import pandas as pd
3 import joblib
4
5 # Load trained model
6 loaded_model = joblib.load("final_attrition_model.pkl")
7
8 # Columns model was trained on
9 model_columns = X_train.columns.tolist()
10
11 # Columns scaler was fitted on
12 scale_cols = list(scaler_hr_attrition.feature_names_in_)    # MUST MATCH YOUR VARIABLE
13
14 def predict_attrition(age, income, overtime, distance, satisfaction):
15
16     # 1. Create a blank row with ALL model columns
17     row = pd.DataFrame([[0] * len(model_columns)], columns=model_columns)
18
19     # 2. Fill user values
20     if "Age" in row.columns: row["Age"] = age
21     if "MonthlyIncome" in row.columns: row["MonthlyIncome"] = income
22     if "DistanceFromHome" in row.columns: row["DistanceFromHome"] = distance
23     if "JobSatisfaction" in row.columns: row["JobSatisfaction"] = satisfaction
24
25     # 3. OverTime dummy
26     if "OverTime_Yes" in row.columns:
27         row["OverTime_Yes"] = 1 if overtime == "Yes" else 0
28
29     # 4. Fill ALL other scaler-required columns with 0 before scaling
30     for col in scale_cols:
31         if col not in row.columns:
32             row[col] = 0
33
34     # 5. Apply scaler using EXACT columns expected during training
35     row[scale_cols] = scaler_hr_attrition.transform(row[scale_cols])
36
37     # 6. Predict
38     pred = loaded_model.predict(row)[0]
39     return "Likely to Leave" if pred == 1 else "Likely to Stay"
40
41
42 # Build the Gradio UI
43 interface = gr.Interface(
44     fn=predict_attrition,
45     inputs=[
46         gr.Number(label="Age"),
47         gr.Number(label="Monthly Income"),
48         gr.Radio(["Yes", "No"], label="OverTime"),
49         gr.Number(label="Distance From Home (km)"),
50         gr.Slider(1, 4, label="Job Satisfaction")
51     ],
52     outputs="text",
53     title="Employee Attrition Prediction System",
54     description="Enter employee details to predict whether they are likely to leave."
55 )
56
57 interface.launch(share=True, debug=True)
```

Colab notebook detected. This cell will run indefinitely so that you can see errors and logs. To turn off, set debug=False in la
* Running on public URL: <https://865527703a8d4be665.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the wor

Employee Attrition Prediction System

Enter employee details to predict whether they are likely to leave.

<p>Age</p> <input type="text" value="41"/>	<p>output</p> <p>Likely to Stay</p>
<p>Monthly Income</p> <input type="text" value="3000"/>	<p>Flag</p>
<p>OverTime</p> <p><input type="radio"/> Yes <input checked="" type="radio"/> No</p>	
<p>Distance From Home (km)</p> <input type="text" value="45"/>	
<p>Job Satisfaction</p> <input type="text" value="1.51"/>	<input type="button" value="Go"/>

Next steps:  [Deploy to Cloud Run](#)