

```

1 # Task 1: Data Loading & Cleaning
2
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7

```

```

1
2 # Load the Dataset:
3
4 file_path = "mini_lending_club.csv"
5 df = pd.read_csv(file_path)
6
7 # Preview the first rows
8 print("Preview of dataset:")
9 display(df.head())
10
11 print("\nDataset shape:", df.shape)

```

Preview of dataset:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_length	home_ownership	annual_inc	purpose	dti	de
0	16795	36 months	23.39%	55.32	G	A5	2 years	RENT	33704	home_improvement	12.11	
1	1860	60 months	29.96%	664.55	F	C1	6 years	OWN	78933	debt_consolidation	1.50	
2	12284	36 months	28.33%	319.87	E	C1	10+ years	OWN	91101	home_improvement	25.18	
3	7265	36 months	21.06%	455.57	E	B3	10+ years	RENT	63856	credit_card	33.34	
4	17850	36 months	15.53%	557.98	D	A1	8 years	OWN	57256	other	9.24	

Dataset shape: (2000, 16)

```

1 # Inspect Generic Info About the Dataset
2 # -----
3
4 print("\nDataset Info:")
5 df.info()
6
7 # Summary statistics for numerical features
8 print("\nSummary Statistics:")
9 display(df.describe())

```

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_amnt        2000 non-null    int64  
 1   term              2000 non-null    object  
 2   int_rate          1900 non-null    object  
 3   installment       2000 non-null    float64 
 4   grade             2000 non-null    object  
 5   sub_grade         2000 non-null    object  
 6   emp_length        1900 non-null    object  
 7   home_ownership    2000 non-null    object  
 8   annual_inc        2000 non-null    int64  
 9   purpose            1900 non-null    object  
 10  dti               2000 non-null    float64 
 11  delinq_2yrs       2000 non-null    int64  
 12  revol_util        1900 non-null    object  
 13  total_acc         2000 non-null    int64  
 14  earliest_cr_line  2000 non-null    int64  
 15  loan_status        2000 non-null    object  
dtypes: float64(2), int64(5), object(9)
memory usage: 250.1+ KB
```

## Summary Statistics:

	loan_amnt	installment	annual_inc	dti	delinq_2yrs	total_acc	earliest_cr_line	grid icon
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	bar icon
mean	18031.749500	413.815520	85211.457000	17.889410	2.070000	40.250000	2006.874000	
std	9788.029854	213.172169	37139.980987	10.114553	1.441575	22.706041	5.899367	
min	1009.000000	50.180000	20057.000000	0.020000	0.000000	1.000000	1999.000000	
25%	9485.500000	230.990000	52672.250000	9.222500	1.000000	20.000000	2001.000000	
50%	18330.500000	411.110000	85122.500000	17.765000	2.000000	40.000000	2008.000000	
75%	26567.250000	589.030000	117562.500000	26.510000	3.000000	60.000000	2012.000000	
max	34997.000000	799.180000	149986.000000	35.000000	4.000000	79.000000	2016.000000	

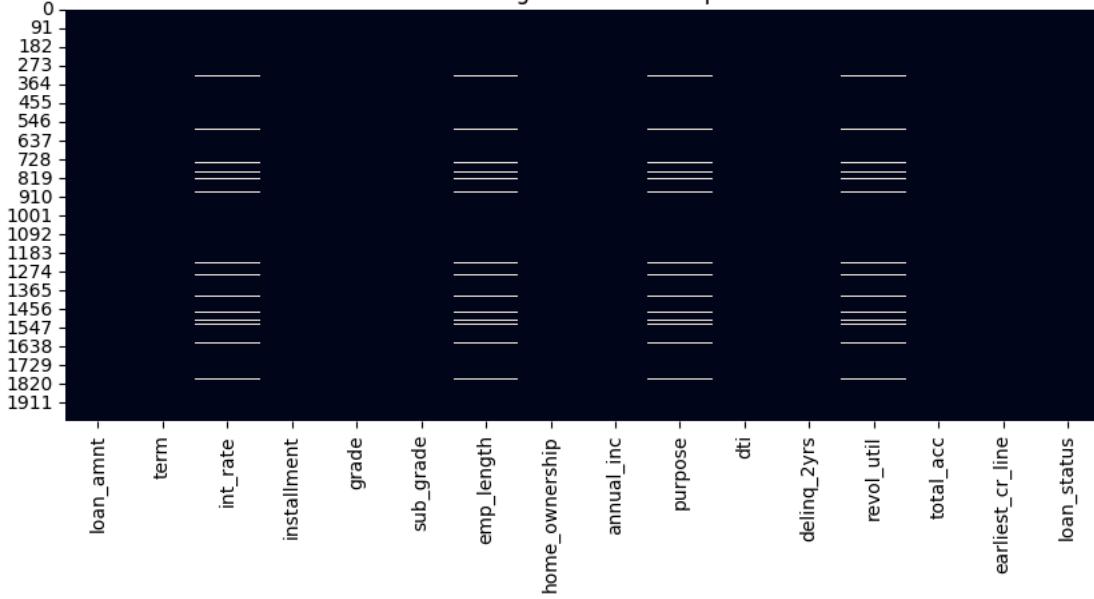
```
1 # 3. Check for missing values
2 # -----
3 print("\nMissing Values:")
4 display(df.isnull().sum())
5
6 # Visualize missing values
7 plt.figure(figsize=(10,4))
8 sns.heatmap(df.isnull(), cbar=False)
9 plt.title("Missing Values Heatmap")
10 plt.show()
11
```

Missing Values:

	0
loan_amnt	0
term	0
int_rate	100
installment	0
grade	0
sub_grade	0
emp_length	100
home_ownership	0
annual_inc	0
purpose	100
dti	0
delinq_2yrs	0
revol_util	100
total_acc	0
earliest_cr_line	0
loan_status	0

dtype: int64

Missing Values Heatmap



```

1 # 4. Clean & preprocess raw string features (SAFE VERSION)
2 # -----
3
4 # Ensure string conversion before using .str
5 df['int_rate'] = df['int_rate'].astype(str).str.replace('%', '', regex=False)
6 df['revol_util'] = df['revol_util'].astype(str).str.replace('%', '', regex=False)
7
8 # Convert to numeric (coerce errors turns bad values into NaN)
9 df['int_rate'] = pd.to_numeric(df['int_rate'], errors='coerce')
10 df['revol_util'] = pd.to_numeric(df['revol_util'], errors='coerce')
11
12 # emp_length cleaning
13 df['emp_length'] = df['emp_length'].astype(str)
14
15 df['emp_length'] = (
16     df['emp_length']
17     .replace({
18         '< 1 year': '0',

```

```

19     '10+ years': '10'
20   })
21   .str.extract(r'(\d+)', expand=False) # extract number
22 )
23
24 df['emp_length'] = pd.to_numeric(df['emp_length'], errors='coerce')
25
26 # earliest credit line
27 df['earliest_cr_line'] = pd.to_numeric(df['earliest_cr_line'], errors='coerce')
28
29
30 # Handle missing values
31 num_cols = df.select_dtypes(include=['float64', 'int64']).columns
32 df[num_cols] = df[num_cols].fillna(df[num_cols].median())
33
34 cat_cols = df.select_dtypes(include=['object']).columns
35 df[cat_cols] = df[cat_cols].fillna(df[cat_cols].mode().iloc[0])
36
37 # -----
38 # 6. Convert loan_status to a binary target
39 #
40
41 df['loan_default'] = df['loan_status'].apply(
42     lambda x: 1 if x == "Charged Off" else 0
43 )
44
45 print("\nLoan Status Conversion (value counts):")
46 print(df['loan_default'].value_counts())
47
48

```

Loan Status Conversion (value counts):

```

loan_default
0    1518
1     482
Name: count, dtype: int64

```

```

1 df.to_parquet("df_cleaned.parquet", index=False)
2

```

```

1 #Task 2 - FEATURE ENGINEERING
2 #
3
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.preprocessing import OneHotEncoder
6 from sklearn.compose import ColumnTransformer
7 from sklearn.pipeline import Pipeline
8
9
10 # These help the model understand borrower risk better.
11
12 df["income_to_loan"] = df["annual_inc"] / df["loan_amnt"]           # Ability to repay
13 df["credit_util_ratio"] = df["revol_util"] / 100                      # Revolving credit usage
14 df["dti_income_ratio"] = df["dti"] / (df["annual_inc"] + 1)          # DTI relative to income
15 df["credit_history_years"] = 2024 - df["earliest_cr_line"]            # Borrower's credit experience
16 df["installment_ratio"] = df["installment"] / (df["annual_inc"] + 1)  # Payment pressure

```

```

1 # SEPARATE FEATURES (X) AND TARGET (y)
2 #
3
4 y = df["loan_default"] # Target is the dependent variable. what I want to predict and it can't be a part of the data that
5
6 # Remove columns that should not be used in modeling, it could lead to data-leakage.
7 drop_cols = ["loan_status", "loan_default", "earliest_cr_line"]
8 X = df.drop(columns=drop_cols)
9
10 # IDENTIFY NUMERIC & CATEGORICAL COLUMNS
11 #
12 numeric_features = X.select_dtypes(include=["float64", "int64"]).columns
13 categorical_features = X.select_dtypes(include=["object"]).columns
14
15 print("Numeric Columns:\n", list(numeric_features))
16 print("\nCategorical Columns:\n", list(categorical_features))
17

```

```

18 y.to_frame().to_parquet("y.parquet", index=False) # save target variable y
19

Numeric Columns:
['loan_amnt', 'int_rate', 'installment', 'emp_length', 'annual_inc', 'dti', 'delinq_2yrs', 'revol_util', 'total_acc', 'income_t']

Categorical Columns:
['term', 'grade', 'sub_grade', 'home_ownership', 'purpose']

1 df.columns
2

Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
       'emp_length', 'home_ownership', 'annual_inc', 'purpose', 'dti',
       'delinq_2yrs', 'revol_util', 'total_acc', 'earliest_cr_line',
       'loan_status', 'income_to_loan', 'credit_util_ratio',
       'dti_income_ratio', 'credit_history_years', 'installment_ratio',
       'loan_default'],
      dtype='object')

1 # SET UP TRANSFORMERS
2 # -----
3
4 # StandardScaler:
5 # → Makes numeric features comparable by converting them to mean=0, variance=1.
6
7 numeric_transformer = StandardScaler()
8
9 # OneHotEncoder:
10 # → Converts categories (e.g., "RENT", "OWN") into 0/1 columns the model can understand.
11 categorical_transformer = OneHotEncoder(handle_unknown='ignore')
12
13 # ColumnTransformer:
14 # → Applies transformations to different column groups:
15 #     - scale numeric features
16 #     - one-hot encode categorical features
17 # → Combines everything into ONE clean matrix
18 # In plain sense I am using the in-built libr ColumnTransf to consolidate both col groups (numerical and Category into one
19
20 preprocessor = ColumnTransformer(
21     transformers=[
22         ("num", numeric_transformer, numeric_features), # scale numerics
23         ("cat", categorical_transformer, categorical_features) # encode categoricals
24     ]
25 )
26
27 # -----
28 # APPLY TRANSFORMATIONS
29 #
30 # fit_transform:
31 # → Learns how to scale from data
32 # → Creates a transformed feature matrix ready for modelling
33
34 X_prepared = preprocessor.fit_transform(X)
35
36 # =====
37 # FIX FEATURE NAMES (IMPORTANT FOR SHAP)
38 # =====
39
40 # 1. Numeric feature names
41 numeric_feature_names = numeric_features.tolist()
42
43 # 2. One-hot encoded categorical feature names
44 categorical_feature_names = (
45     preprocessor.named_transformers_['cat']
46     .get_feature_names_out(categorical_features)
47     .tolist()
48 )
49
50 # 3. Combine them into one final list
51 all_feature_names = numeric_feature_names + categorical_feature_names
52
53 print("Number of feature names:", len(all_feature_names)) # the len indicate the number of features
54
55 import scipy.sparse # for saving the file just in case it crashes
56 scipy.sparse.save_npz("X_prepared.npz", X_prepared)

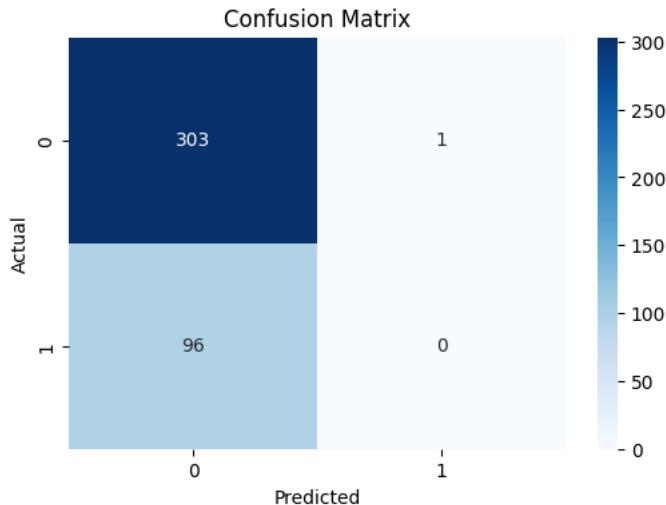
```

Number of feature names: 66

```
1 # Task 3: MODEL TRAINING & EVALUATION
2 # =====
3
4 # Import Suitable Libraries
5
6 from sklearn.model_selection import train_test_split
7 from sklearn.ensemble import RandomForestClassifier
8 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
9
10 # 1. TRAIN-TEST SPLIT
11 # -----
12
13 # X_prepared = transformed feature matrix from Task 2 (pipeline output)
14 # y = target variable from Task 2
15
16 X_train, X_test, y_train, y_test = train_test_split(
17     X_prepared, # this replaces the X value for training
18     y, # this represents the target value
19     test_size=0.2,
20     random_state=42,
21     stratify=y # ensures balanced target distribution
22 )
23
24
25 # TRAIN RANDOM FOREST MODEL
26 # -----
27
28 # Model Selection: I am using RandomForest to determine the prediction for the loan_default prediction
29
30 random_forex = RandomForestClassifier(
31     n_estimators=200,
32     max_depth=None,
33     random_state=42
34 )
35
36 random_forex .fit(X_train, y_train)
37
38 # 3. MAKE PREDICTIONS
39 # -----
40
41 y_pred = random_forex.predict(X_test) # - this is the prediction of the model
42
43
44 # 4. MODEL EVALUATION - this is the process of evaluating that the model has a higher accuracy rate or not.
45 # -----
46
47 print("Accuracy:", accuracy_score(y_test, y_pred))
48 print("\nClassification Report:\n", classification_report(y_test, y_pred))
49
50 # Confusion Matrix
51 cm = confusion_matrix(y_test, y_pred)
52 plt.figure(figsize=(6,4))
53 sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
54 plt.title("Confusion Matrix")
55 plt.xlabel("Predicted")
56 plt.ylabel("Actual")
57 plt.show()
58
```

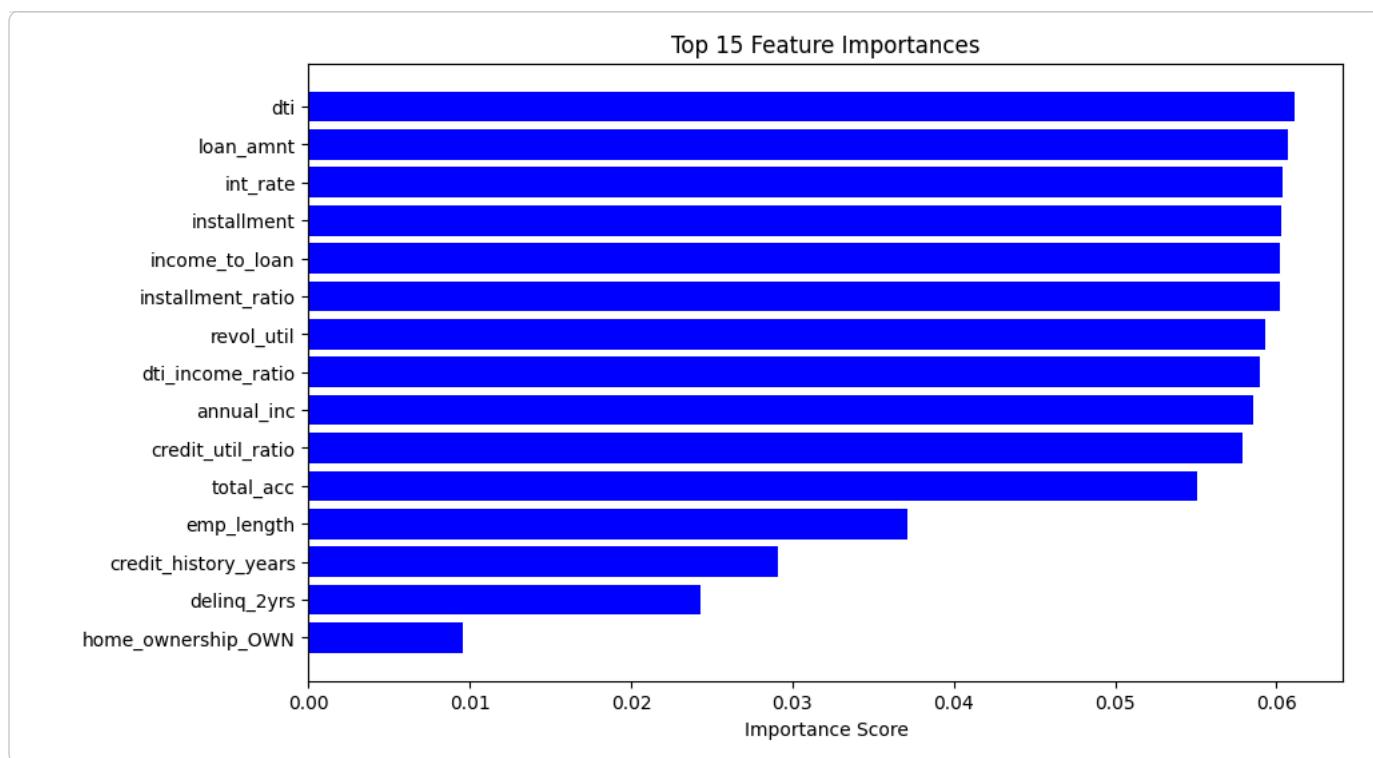
Accuracy: 0.7575

Classification Report:				
	precision	recall	f1-score	support
0	0.76	1.00	0.86	304
1	0.00	0.00	0.00	96
accuracy			0.76	400
macro avg	0.38	0.50	0.43	400
weighted avg	0.58	0.76	0.66	400



```
1 y.to_frame().to_parquet("y.parquet", index=False)
2
```

```
1 # 5. FEATURE IMPORTANCE VISUALIZATION
2 # -----
3
4 # Grab feature names from the preprocessor
5 # (numerical + encoded categorical names)
6 num_features = numeric_features
7 cat_features = list(preprocessor.named_transformers_['cat'].get_feature_names_out(categorical_features))
8
9 all_feature_names = list(num_features) + cat_features
10
11 importances = random_forest.feature_importances_
12
13 # Sort for visualization
14 indices = np.argsort(importances)[-15:] # top 15 features
15
16 plt.figure(figsize=(10,6))
17 plt.barh(range(len(indices)), importances[indices], color='blue')
18 plt.yticks(range(len(indices)), [all_feature_names[i] for i in indices])
19 plt.title("Top 15 Feature Importances")
20 plt.xlabel("Importance Score")
21 plt.show()
```



```

1 # =====
2 # TASK 4 - SHAP & FAIRNESS ANALYSIS
3 # =====
4 import shap
5 shap.initjs()
6
7 # -----
8 # 1. PREP SHAP INPUT (convert sparse → dense)
9 # I had an issue with this section and request the assistance of Google Gemini
10 # -----
11
12 X_test_sample = X_test[:300] # smaller sample for speed
13
14 # Convert sparse matrix to dense
15 if hasattr(X_test_sample, "toarray"):
16     X_test_sample_dense = X_test_sample.toarray()
17 else:
18     X_test_sample_dense = np.array(X_test_sample)
19
20 # -----
21 # 2. SHAP EXPLAINER FOR RANDOM FOREST
22 # -----
23
24 # The error indicates a mismatch in the number of features between shap_values
25 # and X_test_sample_dense. This often happens if the explainer is not
26 # perfectly aligned with the data it's explaining, or if the model's
27 # internal feature representation differs.
28 # To try and align them, we explicitly pass the data to the explainer.
29
30 explainer = shap.TreeExplainer(random_forest, data=X_test_sample_dense)
31
32 # RandomForest produces 2 classes → index 1 is "default"
33
34 shap_values = explainer.shap_values(X_test_sample_dense)
35
36 # Diagnostic: Print shapes before plotting to confirm the issue
37 print(f"Shape of X_test_sample_dense: {X_test_sample_dense.shape}")
38
39 # shap_values can be a list of arrays or a 3D array; handle both
40
41 if isinstance(shap_values, list):
42     print(f"Shape of shap_values (list of arrays): {[s.shape for s in shap_values]}")
43     if len(shap_values) > 1:
44         print(f"Shape of shap_values[1]: {shap_values[1].shape}")
45 else: # assuming it's a numpy array
46     print(f"Shape of shap_values (3D array): {shap_values.shape}")

```

```

47     # The previous line shap_values[1] was incorrect for a 3D array; it picked a sample, not a class.
48     # The correct way to get class 1's shap values for all samples is shap_values[:, :, 1]
49     print(f"Shape of shap_values[:, :, 1]: {shap_values[:, :, 1].shape}")
50 print(f"Length of all_feature_names: {len(all_feature_names)}")
51
52
53 # -----
54 # 3. GLOBAL SHAP SUMMARY PLOTS
55 # -----
56
57 # Beeswarm plot – feature impact
58 shap.summary_plot(
59     shap_values[:, :, 1],           # Correctly select SHAP values for class 1 (default) across all samples
60     X_test_sample_dense,
61     feature_names=all_feature_names
62 )
63
64 # Clean bar plot – easy for reports
65 shap.summary_plot(
66     shap_values[:, :, 1],
67     X_test_sample_dense,
68     feature_names=all_feature_names,
69     plot_type="bar"
70 )
71
72 # -----
73 # 4. LOCAL EXPLANATION (single borrower)
74 # -----
75
76 idx = 0
77 shap.force_plot(
78     explainer.expected_value[1],
79     shap_values[idx, :, 1], # Correctly select SHAP values for a single sample (idx) and class 1
80     X_test_sample_dense[idx],
81     feature_names=all_feature_names
82 )
83
84 # -----
85 # 5. FAIRNESS ANALYSIS – GROUP ACCURACY
86 # -----
87
88 # Recreate original X_test (non-transformed) for grouping
89 from sklearn.model_selection import train_test_split
90
91 X_train_df, X_test_df, y_train_df, y_test_df = train_test_split(
92     X, y, test_size=0.2, random_state=42, stratify=y
93 )
94
95 # Attach predictions to original test data
96 fairness_df = X_test_df.copy()
97 fairness_df["y_true"] = y_test_df.values
98 fairness_df["y_pred"] = random_forest.predict(X_test)
99
100 # Choose grouping variable
101 group_col = "home_ownership"
102
103 print("Groups:", fairness_df[group_col].unique())
104
105 group_metrics = {}
106 for group in fairness_df[group_col].unique():
107     sub = fairness_df[fairness_df[group_col] == group]
108     acc = (sub["y_true"] == sub["y_pred"]).mean()
109     group_metrics[group] = acc
110     print(f"Accuracy for {group}: {acc:.3f}")
111
112 # Bar chart for fairness visualization
113 plt.figure(figsize=(6,4))
114 plt.bar(group_metrics.keys(), group_metrics.values(), color="steelblue")
115 plt.title("Model Accuracy by Home Ownership Group")
116 plt.xlabel("Group")
117 plt.ylabel("Accuracy")
118 plt.ylim(0, 1)
119 plt.show()

```

