

```

1 # Import various libraries
2
3 import pandas as pd
4 import numpy as np
5
6 # Load the dataset:
7
8 hist_stock_prices = pd.read_csv('historical_stock_prices.csv', on_bad_lines='skip')
9 hist_stocks = pd.read_csv('historical_stocks.csv')
10
11 print(hist_stock_prices.head())
12 print("-----")
13 print(hist_stocks.head())

```

	ticker	open	close	adj_close	low	high	volume	date
0	AHH	11.50	11.58	8.493155	11.25	11.68	4633900.0	2013-05-08
1	AHH	11.66	11.55	8.471151	11.50	11.66	275800.0	2013-05-09
2	AHH	11.55	11.60	8.507822	11.50	11.60	277100.0	2013-05-10
3	AHH	11.63	11.65	8.544494	11.55	11.65	147400.0	2013-05-13
4	AHH	11.60	11.53	8.456484	11.50	11.60	184100.0	2013-05-14

	ticker	exchange	name	sector
0	PIH	NASDAQ	PROPERTY INSURANCE HOLDINGS, INC.	FINANCE
1	PIHPP	NASDAQ	PROPERTY INSURANCE HOLDINGS, INC.	FINANCE
2	TURN	NASDAQ	180 DEGREE CAPITAL CORP.	FINANCE
3	FLWS	NASDAQ	1-800 FLOWERS.COM, INC.	CONSUMER SERVICES
4	FCCY	NASDAQ	1ST CONSTITUTION BANCORP (NJ)	FINANCE

	industry
0	PROPERTY-CASUALTY INSURERS
1	PROPERTY-CASUALTY INSURERS
2	FINANCE/INVESTORS SERVICES
3	OTHER SPECIALTY STORES
4	SAVINGS INSTITUTIONS

```

1 # Data Inspection:
2
3 print(hist_stock_prices.info())
4 print("-----")
5 print(hist_stocks.info())
6
7 print("-----")
8 print("Summary---Stats")
9 print(hist_stock_prices.describe(), hist_stocks.describe())
10
11

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1747016 entries, 0 to 1747015
Data columns (total 8 columns):
 #   Column      Dtype  
 --- 
 0   ticker      object 
 1   open        float64
 2   close       float64
 3   adj_close   float64
 4   low         float64
 5   high        float64
 6   volume      float64
 7   date         object 
dtypes: float64(6), object(2)
memory usage: 106.6+ MB
None
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6460 entries, 0 to 6459
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   ticker      6460 non-null   object 
 1   exchange    6460 non-null   object 
 2   name        6460 non-null   object 
 3   sector      5020 non-null   object 
 4   industry    5020 non-null   object 
dtypes: object(5)
memory usage: 252.5+ KB
None
-----
Summary---Stats

```

```

      open      close     adj_close      low      high \
count  1.747016e+06  1.747016e+06  1.747016e+06  1.747015e+06
mean   2.314312e+02  2.317348e+02  2.272249e+02  2.265874e+02  2.365962e+02
std    7.573427e+03  7.604517e+03  7.603762e+03  7.431390e+03  7.747283e+03
min   5.000000e-03  5.000000e-03  1.464524e-05  5.000000e-03  5.000000e-03
25%   6.810000e+00  6.810487e+00  4.378336e+00  6.680000e+00  6.938271e+00
50%   1.471857e+01  1.471857e+01  1.109731e+01  1.450000e+01  1.493750e+01
75%   2.975000e+01  2.975000e+01  2.509826e+01  2.933333e+01  3.013000e+01
max   8.028744e+05  8.002500e+05  8.002500e+05  7.290000e+05  8.190000e+05

      volume
count  1.747015e+06
mean   2.885017e+06
std    4.245034e+07
min   1.000000e+00
25%   2.190000e+04
50%   1.231000e+05
75%   5.387000e+05
max   4.483504e+09      ticker exchange
count   6460      6460      6460      5020
unique   6460          2      5462      13
top      ZYME      NASDAQ  BANK OF AMERICA CORPORATION  FINANCE
freq       1      3308          16      1022

industry

```

Strategies I applied to the Data Cleaning for the Stock Prices File:

The raw historical stock price dataset contained multiple extreme and unrealistic values, including adjusted closing prices in the 10^{19} range and stock prices exceeding two million dollars. These outliers were identified using summary statistics and removed by applying reasonable bounds for valid stock price ranges (0.1–10,000) and trading volumes (0–500 million shares). Missing or malformed dates were cleaned by converting the date column to datetime format and dropping invalid entries. Duplicate rows (ticker–date combinations) were removed to ensure accurate time-series analysis. After cleaning, all price and volume columns contained realistic values, and the dataset was safely prepared for analysis."

```

1 # Data Cleaning & Processing
2
3 # Convert Date to Date_Time:
4
5 hist_stock_prices['date'] = pd.to_datetime(hist_stock_prices['date'], errors='co
6 hist_stock_prices = hist_stock_prices.dropna(subset=['date'])
7
8 # Drop the Duplicates for Ticker and Date: This ensures that the ticker and date
9
10 hist_stock_prices = hist_stock_prices.drop_duplicates(subset=['ticker', 'date'])
11
12 # Fix Extreme Price Outliers in the dataset and remove unrealistic Prices (Outlie
13
14 valid_prices = (hist_stock_prices['close'] > 0.1) & (hist_stock_prices['close'] <
15 hist_stock_prices = hist_stock_prices[valid_prices]
16
17 for col in ['open', 'high', 'low', 'adj_close']: # apply the same range to the o
18     valid_range = (hist_stock_prices[col] > 0.1) & (hist_stock_prices[col] < 1000)
19     hist_stock_prices = hist_stock_prices[valid_range]
20
21 # Handle the volume col based on the reality:
22
23 valid_volume = (hist_stock_prices['volume'] > 0) & (hist_stock_prices['volume'] <
24 hist_stock_prices = hist_stock_prices[valid_volume]
25
26 # Remove Rows that have missing values to ensure that the price cols have usable
27
28 price_cols=['open', 'high', 'low', 'close', 'adj_close'] # ensures missing rows a
29 hist_stock_prices = hist_stock_prices.dropna(subset=price_cols)
30
31 # Sort by Ticker (Unique Identifier) & the Date:
32
33 hist_stock_prices = hist_stock_prices.sort_values(by=['ticker', 'date'])
34
35 # Reset the Index - this syntax ensures that the index for the dataset is reset a
36 hist_stock_prices = hist_stock_prices.reset_index(drop=True)
37
38 hist_stock_prices = hist_stock_prices.set_index('date')

```

```

39
40 # Check the execution:
41
42 # print(hist_stock_prices.describe())
43 print(hist_stock_prices.info(), hist_stock_prices.head())
44

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1714712 entries, 1987-09-23 to 2018-08-24
Data columns (total 7 columns):
 #   Column      Dtype  
--- 
 0   ticker      object  
 1   open         float64 
 2   close        float64 
 3   adj_close    float64 
 4   low          float64 
 5   high         float64 
 6   volume       float64 
dtypes: float64(6), object(1)
memory usage: 104.7+ MB
None           ticker      open     close   adj_close      low     high   \
date
1987-09-23    AAPL     1.933036  1.973214  0.101052  1.919643  2.000000
1987-09-24    AAPL     1.973214  2.017857  0.103339  1.973214  2.066964
1987-09-25    AAPL     2.026786  2.053571  0.105168  2.017857  2.071429
1987-09-28    AAPL     2.053571  1.991071  0.101967  1.982143  2.098214
1987-09-30    AAPL     1.937500  2.017857  0.103339  1.937500  2.035714

volume
date
1987-09-23  63644000.0
1987-09-24  45640000.0
1987-09-25  26630800.0
1987-09-28  50960000.0
1987-09-30  30520000.0

```

```

1 # Checking Cols in Historical Stock Dataset:
2
3 print(hist_stocks.info())
4 print(hist_stocks.describe(include='all'))

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6460 entries, 0 to 6459
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   ticker      6460 non-null   object  
 1   exchange    6460 non-null   object  
 2   name        6460 non-null   object  
 3   sector      5020 non-null   object  
 4   industry    5020 non-null   object  
dtypes: object(5)
memory usage: 252.5+ KB
None           ticker      exchange      name      sector   \
count    6460      6460            6460      5020
unique   6460      2              5462      13
top      ZYME      NASDAQ      BANK OF AMERICA CORPORATION FINANCE
freq     1        3308            16      1022

industry
count      5020
unique    136
top      MAJOR PHARMACEUTICALS
freq     419

```

```

1 # Create the Decade Cols ( Cols Segmentation)
2 #The best practice would be to create a decade col before merging to avoid comput
3
4 # Extract the year
5 hist_stock_prices['year'] = hist_stock_prices.index.year # extracting the year co
6
7 # Convert Year into Decades
8 hist_stock_prices['decade'] = (hist_stock_prices['year'] // 10) * 10 # this syntax
9
10 hist_stock_prices['decade'] = hist_stock_prices['decade'].astype(str) + "s" # cha
11
12 # Segment Decade Cols
13 decades = sorted(hist_stock_prices['decade'].unique())
14 print(decades)

```

```

15
16 # Construct a Dataframe using the Dictionary method:
17
18 decades_dfs ={decade: hist_stock_prices[hist_stock_prices['decade'] == decade]
19             for decade in decades}
20
21 print(decades_dfs["1990s"].head())
22

['1970s', '1980s', '1990s', '2000s', '2010s']
      ticker    open     close   adj_close      low      high \
date
1990-01-02  AAPL  1.258929  1.330357  0.122946  1.250000  1.339286
1990-01-03  AAPL  1.357143  1.339286  0.123771  1.339286  1.357143
1990-01-04  AAPL  1.366071  1.343750  0.124184  1.330357  1.383929
1990-01-05  AAPL  1.348214  1.348214  0.124597  1.321429  1.366071
1990-01-08  AAPL  1.339286  1.357143  0.125422  1.321429  1.357143

      volume   year decade
date
1990-01-02  45799600.0  1990  1990s
1990-01-03  51998800.0  1990  1990s
1990-01-04  55378400.0  1990  1990s
1990-01-05  30828000.0  1990  1990s
1990-01-08  25393200.0  1990  1990s

```

The metadata cleaning process successfully standardized all text-based columns, removed extra whitespace, handled missing values, and optimized categorical variables. The 'ticker', 'exchange', 'sector', and 'industry' columns were converted to categorical data types, improving memory efficiency and enabling faster analytical operations. Missing sector and industry values were filled with a placeholder category ("Unknown"), ensuring no null values remain in any categorical fields. After cleaning, the dataset retains all 6,460 unique companies, and the previews confirm that fields are well formatted and ready for merging with the stock price dataset.

```

1 # Data Cleaning Process
2
3 # Re-load hist_stocks to ensure it's a DataFrame
4 hist_stocks = pd.read_csv('historical_stocks.csv')
5
6 # REMOVE DUPLICATE ROWS BASED ON 'ticker' this is important before merging with t
7 hist_stocks = hist_stocks.drop_duplicates(subset=['ticker'])
8
9 # CLEAN TEXT COLUMNS (remove extra spaces, standardize case)
10
11 text_cols=['ticker','exchange','name','sector','industry']
12
13 for col in text_cols:
14     hist_stocks[col] = hist_stocks[col].astype(str).str.strip() # this syntax remov
15     hist_stocks[col] = hist_stocks[col].replace('NaN', None) # this replaces NaN an
16
17 # FILL MISSING VALUES FOR SECTOR & INDUSTRY:
18
19 hist_stocks['sector'] = hist_stocks['sector'].fillna('Unknown')
20 hist_stocks['industry'] = hist_stocks['industry'].fillna('Unknown')
21
22 # Convert Cols to Category:
23
24 categorical_cols = ['exchange', 'sector', 'industry'] # this is a list that used
25
26 for cols in categorical_cols:
27     hist_stocks[cols] = hist_stocks[cols].astype('category')
28
29 hist_stocks = hist_stocks.reset_index(drop=True)
30
31 # Check execution:
32 print("Inspect the Dataset & First 5 Heads")
33 print(hist_stocks.info())
34 print(hist_stocks.head())
35

```

```

Inspect the Dataset & First 5 Heads
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6460 entries, 0 to 6459
Data columns (total 5 columns):

```

```
# Column Non-Null Count Dtype
--- -----
0 ticker    6460 non-null   object
1 exchange  6460 non-null   category
2 name      6460 non-null   object
3 sector    6460 non-null   category
4 industry  6460 non-null   category
dtypes: category(3), object(2)
memory usage: 132.2+ KB
None
    ticker exchange           name          sector \
0 PIH     NASDAQ  PROPERTY INSURANCE HOLDINGS, INC. FINANCE
1 PIHPP   NASDAQ  PROPERTY INSURANCE HOLDINGS, INC. FINANCE
2 TURN    NASDAQ  180 DEGREE CAPITAL CORP. FINANCE
3 FLWS    NASDAQ  1-800 FLOWERS.COM, INC. CONSUMER SERVICES
4 FCCY    NASDAQ  1ST CONSTITUTION BANCORP (NJ) FINANCE

industry
0 PROPERTY-CASUALTY INSURERS
1 PROPERTY-CASUALTY INSURERS
2 FINANCE/INVESTORS SERVICES
3 OTHER SPECIALTY STORES
4 SAVINGS INSTITUTIONS
```

```
1 # Merging the Dataset into Dataframe ( using the most common value - ticker (lef
2
3 # Reset_Index so Date Col become an Index for merging:
4
5 hist_stock_prices = hist_stock_prices.reset_index()
6
7 merge_decade_df = hist_stock_prices.merge(hist_stocks, on='ticker', how='left')
8
9 merged_decade_df = merge_decade_df.set_index('date')
10
11 print(merged_decade_df.info())
12 print(merged_decade_df.head())
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1714712 entries, 1987-09-23 to 2018-08-24
Data columns (total 13 columns):
 # Column Dtype
--- -----
0 ticker    object
1 open      float64
2 close     float64
3 adj_close float64
4 low       float64
5 high      float64
6 volume    float64
7 year      int32
8 decade    object
9 exchange  category
10 name     object
11 sector    category
12 industry  category
dtypes: category(3), float64(6), int32(1), object(3)
memory usage: 143.9+ MB
None
    ticker      open      close      adj_close      low      high   \
date
1987-09-23  AAPL  1.933036  1.973214  0.101052  1.919643  2.000000
1987-09-24  AAPL  1.973214  2.017857  0.103339  1.973214  2.066964
1987-09-25  AAPL  2.026786  2.053571  0.105168  2.017857  2.071429
1987-09-28  AAPL  2.053571  1.991071  0.101967  1.982143  2.098214
1987-09-30  AAPL  1.937500  2.017857  0.103339  1.937500  2.035714

    volume      year      decade      exchange      name      sector   \
date
1987-09-23  63644000.0  1987  1980s  NASDAQ  APPLE INC. TECHNOLOGY
1987-09-24  45640000.0  1987  1980s  NASDAQ  APPLE INC. TECHNOLOGY
1987-09-25  26630800.0  1987  1980s  NASDAQ  APPLE INC. TECHNOLOGY
1987-09-28  50960000.0  1987  1980s  NASDAQ  APPLE INC. TECHNOLOGY
1987-09-30  30520000.0  1987  1980s  NASDAQ  APPLE INC. TECHNOLOGY

    industry
date
1987-09-23  COMPUTER MANUFACTURING
1987-09-24  COMPUTER MANUFACTURING
1987-09-25  COMPUTER MANUFACTURING
1987-09-28  COMPUTER MANUFACTURING
1987-09-30  COMPUTER MANUFACTURING
```

"The summary statistics reveal a steady increase in average stock prices across decades. Earlier decades such as the 1970s and 1980s show lower mean and median values, while the 1990s and 2000s present higher averages, indicating market growth. Standard deviation increases in the later decades, suggesting greater volatility and larger price movements during periods such as the dot-com boom in the late 1990s and the financial crisis of the late 2000s."

```

1 # Summary Stats of the Dataset (finding the mean, median, max and std for each decade)
2
3 summary_stats = {}
4
5 for decade, df in decades_dfs.items():
6     summary_stats[decade] = df[['open', 'high', 'low', 'close', 'volume']].agg(['mean', 'median', 'max', 'std'])
7
8 for decade, stats in summary_stats.items():
9     print(f"Summary Statistics for {decade}:")
10    print("-----")
11    print(stats)
12

```

Summary Statistics for 1970s:

	open	high	low	close	volume
mean	4.905704	4.958258	4.857528	4.906648	3.304594e+05
median	6.265625	6.312500	6.218750	6.265625	2.288000e+05
max	12.875000	13.750000	12.875000	13.750000	8.934400e+06
std	3.664308	3.694529	3.630428	3.666292	3.690889e+05

Summary Statistics for 1980s:

	open	high	low	close	volume
mean	9.293568	9.422624	9.186185	9.293259	4.491279e+05
median	5.145833	5.250000	5.093750	5.140506	5.360000e+04
max	187.101242	188.786835	186.679840	187.733337	2.563540e+08
std	14.647354	14.837237	14.444184	14.650037	2.906768e+06

Summary Statistics for 1990s:

	open	high	low	close	volume
mean	21.704312	22.075988	21.326369	21.702580	1.449757e+06
median	9.875000	10.000000	9.750000	9.875000	6.800000e+04
max	989.777771	997.777771	980.000000	996.888916	4.747008e+08
std	58.594048	59.746380	57.368122	58.531888	9.774339e+06

Summary Statistics for 2000s:

	open	high	low	close	volume
mean	26.907055	27.406056	26.369693	26.890713	2.734352e+06
median	15.280000	15.531250	15.015000	15.280000	1.439000e+05
max	998.400024	999.799988	996.799988	998.500000	4.959490e+08
std	53.760724	54.969744	52.379907	53.632287	1.748728e+07

Summary Statistics for 2010s:

	open	high	low	close	volume
mean	35.193171	35.623808	34.747779	35.191381	1.366161e+06
median	20.270000	20.549999	20.000000	20.274745	1.494000e+05
max	999.000000	999.000000	986.729980	992.039978	4.994696e+08
std	58.989201	59.717626	58.242893	58.972486	8.134466e+06

The monthly average close price plots display an upward long-term trend across all decades, reflecting overall market growth. The 1990s and early 2000s show pronounced fluctuations, consistent with the dot-com era and increased market speculation. The 2000s also exhibit a significant decline around 2008, corresponding to the global financial crisis. The smoother movement in earlier decades indicates lower volatility, while later decades show more dynamic price behavior due to increased market participation and technological changes."

```



1 # Visualization
2
3 #Monthly average Close price plots for each decade
4
5 import matplotlib.pyplot as plt
6
7 for decade, df in decades_dfs.items():
8     monthly_close = df['close'].resample('M').mean()
9
10    plt.figure(figsize=(10, 5))
11    plt.plot(monthly_close, label=f"{decade} Monthly Avg Close")
12    plt.title(f"Average Monthly Close Prices - {decade}")

```

```
13 plt.xlabel("Year")
14 plt.ylabel("Close Price")
15 plt.grid(True)
16 plt.legend()
17 plt.show()
18
```



```
/tmp/ipython-input-1191909551.py:7: FutureWarning: 'M' is deprecated and will be removed
monthly_close = df['close'].resample('M').mean()
```




James Kiawu
12:47 PM Today





This represents the reference to my sources

```
1 # Vizualization
2
3 # Volume Distribution Histogram for Each Decade
4
5 for decade, df in decades_dfs.items():
6     plt.figure(figsize=(10, 5))
7     plt.hist(df['volume'], bins=50, alpha=0.7)
8     plt.title(f'Volume Distribution - {decade}')
9     plt.xlabel('Volume')
10    plt.ylabel('Frequency')
11    plt.grid(True)
12    plt.show()
13
```







Boxplots for each decade reveal that the price range widens significantly in the later decades, especially the 1990s and 2000s, reflecting increased market volatility. Earlier decades exhibit narrower boxes, indicating more stable price fluctuations. The presence of outliers becomes more pronounced in the 2000s and 2010s, highlighting unusual price spikes or declines likely connected to economic crises, rapid sectoral growth, or major geopolitical events.

```

1 # Vizualization
2
3 # Boxplots for High and Low prices per decade
4
5 for decade, df in decades_dfs.items():
6     plt.figure(figsize=(10, 5))
7     plt.boxplot([df['high'], df['low']], labels=['High', 'Low'])
8     plt.title(f'High & Low Price Distribution - {decade}')
9     plt.ylabel('Price')
10    plt.grid(True)
11    plt.tight_layout()
12    plt.show()

```

